

Master's Proposal

Slava Nikitin

December 9, 2014

Contents

1	Introduction	3
1.1	Formal Modeling Approach	13
1.2	Data-Generating Experimental Tasks	21
1.3	Benchmark Behavioral Data	23
1.4	Theoretical Data Models	30
1.5	Bayesian Statistical Framework	33
1.5.1	Bayesian Hierarchical Models	34
1.5.2	MCMC Background	38
1.5.3	Metropolis-Hastings Algorithm	42
2	Questions to be Answered	47
2.1	Modeling Parameter Dependencies	48
2.2	Predictions of Dependent Parameters	51
2.3	Correlation Structure in a Benchmark Dataset	53
3	Models	55
3.1	Ratcliff Diffusion Model	56

3.2	Copula-based Multivariate Distributions	70
3.3	Proposed Generalizations	83
4	Proposed studies	89
4.1	Study A: Comparison of Models' Predictions	90
4.1.1	Mean Sample Paths	90
4.1.2	Response times and Responses	93
4.1.3	Parameter Settings	97
4.1.4	Results	102
4.1.5	Discussion	104
4.2	Study B: Benchmark Dataset Analyses	105
4.2.1	Outlier Pre-processing	105
4.2.2	Bayesian Models	108
4.2.3	MCMC Sampler	114
4.2.4	Convergence Diagnostics	120
4.2.5	Characterizing Parameter Dependencies	122
4.2.6	Posterior Predictive Model Comparison	124
4.2.7	Predictive Accuracy of Models	126
4.2.8	Results	130
4.2.9	Discussion	130
5	Bibliography	133

Chapter 1

Introduction

Under the same environmental conditions humans, and other phylogenetically younger animals, show a limited number of alternative behaviors. Variability of behavior suggests that evolution of complex perceptual and cognitive functions presents animals with choices for any given situation. Given that behavior is organized successfully in light of choices, their nervous systems must have a mechanism for transforming competing choices into a decision. Thus, understanding how the brain organizes behavior of complex animals can be advanced by understanding decision mechanisms it employs.

Laboratory investigations of decision-making involving humans span a wide range of decision situations. Humans can be asked to make complex, multi-step decisions involved in selecting among health care policies (Peters et al., 2014), gambles (Tversky and Kahneman, 1981), competitive strategies (Camerer, 2003), and to make simple one-shot decisions based on a single dimension of a stimulus such as presence of a visual pattern (Smith, 1995),

relative brightness of stimuli (Ratcliff and Rouder, 1998), or motion direction (Ratcliff and McKoon, 2008). In this thesis, I will concentrate on simple one-step decisions among two alternatives because resulting behavioral data can be obtained under rigorous control (Van Zandt and Townsend, 2013), the decision process is more amenable to mathematical modeling (Coombs, 1983; Lewandowsky and Farrell, 2010), and the uncovered principles may generalize to all forms of decisions (Shadlen and Kiani, 2013).

An example of laboratory tasks I concentrate on is a lexical decision task (Wagenmakers, 2009). On each trial, a participant is shown a string of letters, and asked to quickly and accurately decide whether it is a word or not by pressing one of two buttons on a response box. Thus, the design of an experiment couples participant's decision with a simple movement (Van Zandt and Townsend, 2013), and enables easy collection of responses, response times and physiological measures that can provide substantial clues about the simple decision mechanism (Luce, 1986; Shadlen and Kiani, 2013; Mulder et al., 2014).

The data bearing on psychological theory of simple decisions comes from both neural and behavioral studies. On one hand, the key findings from neural studies involve growth and decay shapes of firing rate functions of neurons interpreted as implementing competing decisions (Gold and Shadlen, 2007). On the other hand, behavioral data contributes response rates and response time distributions, which have had the most effect on theory so far and will be the focus of the thesis (Stone, 1960; Edwards, 1965; Vickers, 1979; Townsend and Ashby, 1983; Luce, 1986; Shadlen and Newsome, 1996;

Ratcliff and Rouder, 1998; Smith and Ratcliff, 2004; Bogacz, 2007; Mulder et al., 2014). The primary features of behavioral data include positively skewed response time distributions, asymmetries between correct and error means, speed-accuracy trade off and sequential relations between consecutive responses and response times (Craigmile et al., 2010; Ratcliff, 2014).

The currently dominant theoretical understanding of simple one-stage decisions is based on analogy with sequential sampling of data from statistics (Wald and Wolfowitz, 1948; Stone, 1960; Vickers, 1979; Bogacz et al., 2006). One version of psychological theory consistent with sequential sampling postulates that on every trial a participant’s brain forms a variable representation of the decision-relevant stimulus feature. To make a decision, participants have a mechanism to sample noisy signals from the representation and accumulate them into separate representations standing for competing choices. A decision is formed when one of the choice representations, called evidence integrators, reaches a preset and stable threshold of evidence. Associated movement execution follows depending on which evidence integrator reached the threshold first.

There are several competing models that implement sequential sampling idea in a mathematical formalism (Stone, 1960; Pike, 1973; Ratcliff, 1978; Usher and McClelland, 2001; Brown and Heathcote, 2008). Assumptions of the models provide a precise description of how various parameters of processing combine to form simple decisions, parameters that are supposedly capturing higher-level properties of a task-evoked neural network (Gold and Shadlen, 2007; Cassey et al., 2014; Mulder et al., 2014; Cole et al., 2014). The

models highlight time to encode evidence, rate of evidence uptake, evidence threshold, initial evidence and time to execute a motor response as important processing parameters in explaining patterns of behavioral data, but make different assumptions about how these parameters combine to generate choice behavior.

One way to categorize all the models is based on three types of assumptions that they make. The first kind of assumptions describe the number, interactions and dynamics of evidence integrators during a trial, potentially expressed as a system of stochastic processes (Smith, 2000). The second kind makes assumptions about other stages of processing involved in a trial, such as encoding of stimulus features and motion response, which frequently amounts to assigning an overall time constant that describes their duration (Ratcliff and Smith, 2004). Lastly, simple decision models incorporate distributional assumptions that describe how parameters vary across trials to capture variability in the processing system (Ratcliff and Smith, 2004; Jones and Dzhafarov, 2014).

Mathematical assumptions of the competing models make different theoretical claims about processing, and testable predictions about responses and response time distributions can be deduced from the assumptions (Vickers, 1979; Coombs, 1983; Luce, 1986). A set of models, widely used for providing process-based interpretation of behavior (Ratcliff and McKoon, 2008; Donkin et al., 2011; Heathcote and Hayes, 2012), have been developed that account, with high accuracy, for a large number of patterns in behavioral data including response time distributions, accuracy, and relations between response

times and accuracy, across a variety of manipulations and paradigms (Usher and McClelland, 2001; Ratcliff and Tuerlinckx, 2002; Brown and Heathcote, 2008). The problems with these models are that they mimic each other statistically with common sample sizes while making different theoretical claims (Van Zandt et al., 2000; Ratcliff and Smith, 2004; Donkin et al., 2011), and they do not yet account for the whole pattern of behavioral data.

One of the ways to push toward understanding of a simple decision process is to incorporate additional theoretical principles into promising models and attempt to account for unexplained data patterns. Looking at the history of the sequential sampling field, one of the largest differences between currently popular models and their predecessors is addition of assumptions about variability of parameters across trials (Laming, 1968; Vickers, 1979; Ratcliff and Smith, 2004). Variability assumptions express the idea that processing is not static across trials within a condition. Starting evidence may include evidence remaining from the previous trial or movement time may be faster if the same stimulus repeats. Expressed in terms of probability distributions, variability in parameters can account for additional data involving asymmetries in relations between speed and accuracy of processing, which the old models missed (Luce, 1986; Ratcliff and Smith, 2004).

While some gains in understanding have already been made from incorporating trial-to-trial variability in parameters into models, I suggest that further gains can be made by better characterizing this sort of variability. Consider that currently popular models differ in assuming what parameters vary and what their distributional forms are (Jones and Dzhafarov, 2014),

but they all commit to two questionable claims: on each trial a parameter is set to a value independent of the previous trials and all parameters are mutually independent (Usher and McClelland, 2001; Ratcliff and Tuerlinckx, 2002; Brown and Heathcote, 2008).

The first assumption is challenged by the always present auto-correlation in response times and responses (Peruggia et al., 2002; Jones et al., 2013). Behavioral data suggests that, say, the threshold of a participant on one trial depends on the threshold on the previous several trials. One line of recent model development incorporated the principle of auto-correlated parameters through additional processes including learning of stimulus probabilities, error-correction and residual evidence (Cho et al., 2002; Wagenmakers et al., 2004; Brown et al., 2008; Gao et al., 2009; Goldfarb et al., 2012). Depending on the model, starting evidence, rate of evidence uptake and evidence thresholds were functions of values on previous trials and parameters of additional processes. Such modifications enabled models to predict a variety of sequential effects and improve understanding of decision process dynamics.

However, the other questionable assumption, the issue of mutual independence of processing parameters, has not been examined. An alternative to independence is that processing parameters vary in systematic ways from trial to trial. Several sources of evidence provide support for statistical dependence among parameters.

One form of evidence comes from model-based analyses of response time distributions under block-wise manipulation of speed-accuracy instructions. Using a couple of prominent sequential sampling models, researchers showed

that the best explanation of accuracy emphasis is based on simultaneous rise in the rate of evidence accumulation and evidence threshold (Vandekerckhove and Tuerlinckx, 2007; Vandekerckhove et al., 2011; Rae et al., 2014), and sometimes also increase in non-decision time (Voss et al., 2004; Zhang and Rowe, 2014). In further support of change in non-decision time, a purely empirical study, using measurements of lateralized readiness potential with electroencephalograph, demonstrated that under speed instructions the time interval between onset of motor processing and motor response decreases (Rinkenauer et al., 2004). Both kinds of studies suggest that a control process that sets a speed-accuracy regime, as manipulated by instructions, can modulate activity in two or more parts of a task-evoked neural network and its effects are picked up by the models’ parameters (Wang, 2008; Hare et al., 2011; Turner et al., 2013a; Cole et al., 2014).

Another block-wise manipulation indicating systematic tradeoffs among parameters is due to practice. Using a sequential sampling model to analyze 10,000 trials of the lexical decision task, researchers found a coordinated adjustment of processing parameters with practice (Dutilh et al., 2009). The nondecision time decreased, rate of evidence uptake increased, evidence threshold decreased and initial evidence became less biased with practice. In a similar model-based study, but based on a motion discrimination task, participants showed decrease in the evidence threshold and increase in the rate of evidence uptake over 4032 trials (Zhang and Rowe, 2014). Overall, practice effects are also consistent with a principle that decision and non-decision parameters change in a dependent manner.

On a trial-by-trial level, the phenomenon of post-error slowing also suggests dependencies among parameters as a feature of sequential decision-making (Rabbitt, 1979). Post-error slowing refers to the phenomenon that, when a participant makes a mistake, response time on the next trial slows down. Explanations of this phenomenon have usually evoked only an increase in the threshold, but a model-based analysis showed that rate of evidence uptake, threshold and nondecision time may all be affected (Dutilh et al., 2013). Dutilh et al. (2013) examined old and young adults with lexical decision and motion discrimination tasks. Old adults showed an increase in evidence threshold, decrease in rate of evidence uptake and increase in non-decision time across both tasks. Young adults showed an increase in non-decision time in the motion discrimination task, but in the lexical decision task they increased evidence thresholds and increased non-decision time.

Another sequential effect that is consistent with dependent parameters is based on effects of prior stimuli on response times. Goldfarb et al. (2012) collected behavioral data for a task requiring discrimination between upper and lower case letter Os. They analyzed mean response times, classified by correct and error response, on the second trial of two-stimulus sequences consisting of all possible combinations of repetitions and alternations. Conditioned on stimulus sequence, participants showed trade-offs between correct and error response times: faster response time on error trials corresponded to slower response times on correct trials, or vice versa. A model with trial-varying initial evidence captured the qualitative pattern, but systematically missed response times for error trials. Given the mechanics of sequential sampling

models, the misses may have been due to not accounting for adjustments in thresholds or rate of information uptake that shift simultaneously with initial evidence.

An account of sequential effects based on simultaneous adjustment in parameters has been proposed by Ratcliff et al. (1999). Participants were given a numerosity task where clusters of asterisks need to be discriminated as large or small. Some participants showed elevated probabilities and speed when repeating the previous response rather than altering the response. Explanation of both the response probability and response time was successful only when initial evidence and rate of evidence of uptake were allowed to adjust simultaneously.

Finally, on a more speculative note, one possible basis of parameter correlations is suggested by studies of control processes, such as task instruction processing, task representation, error monitoring, selective attention, that are found to frequently (on the order of seconds) modulate task-evoked brain network, a distributed and interconnected set of neural populations underlying task performance, in line with task demands and statistical structure of stimuli (Mozer et al., 2007; Jones et al., 2009; Collins and Frank, 2013). The neural basis of control processes has been generally assigned to prefrontal cortex and basal ganglia (van Schouwenburg et al., 2010), regions which have two-way projections to almost all the other structures in the limbic system and neocortex (Miller and Cohen, 2001; Calabresi et al., 2014). Hence, current understanding of neural architecture of the control processes provides an opportunity for simultaneous modulation of activity in all components of

a task-evoked network supporting particular speeded decisions.

Behavioral and neural data presented above suggests, without being conclusive, that there may be correlations between processing parameters underlying simple choice behavior. While other interpretations are possible, I propose to quantify evidence for a simple and falsifiable principle of processing dynamics underlying visually-guided choice behavior: From trial to trial, and across conditions, psychological parameters identified by sequential sampling models vary systematically in a dependent manner. The suggested “dependence principle” is in qualitative contrast with the usually assumed independence of processing parameters (Ratcliff and Smith, 2004).

Given the sharp contrast between the proposed and the conventional principle, the goals motivating the proposed thesis are to generalize one of the currently popular sequential sampling models of decision and study its relation to a model with independent parameters. The proposed approach involves adding a correlation structure that can describe correlations between parameters without making a theoretical claim as to how correlations arise. Using a generalized model several questions can be posed and answered: Do models with and without dependencies lead to different time courses of the decision process? Are behavioral predictions distinguishable between them? What can we extract from data about parameter correlations? How much does predictive accuracy increase when adding a correlation structure? In the end, addressing these questions through theoretical and empirical comparisons of the standard version and a generalized version of a sequential sampling model can give better insight into processing underlying choice be-

havior.

In the rest of this introduction chapter, I will lay out reasons for taking formal modeling approach to test the dependence principle and elaborate the structure of experiments generating behavioral data. As an example of a typical experiment, I will present a benchmark dataset, and important response time and accuracy features that will form the focus of comparisons between models. Then, I will present standard sequential sampling framework and then introduce a method of constructing multivariate distributions that will be useful for developing a generalized model of decision making. Finally, I will present a statistical framework grounded in Bayesian theory that will be used to estimate the correlation structure and compare predictive accuracy of proposed models.

1.1 Formal Modeling Approach

Psychology has as one of its chief aims understanding of the cognitive architecture that generates the rich repertoire of human behavior and is sensitive to many environmental variables (Anderson et al., 2004; Anderson, 2007). More specifically, a psychologist wants to learn about the number, temporal arrangement, properties and interactions of elementary processes that organize behavior. A fundamental obstacle to satisfying psychologist’s aspirations is that cognition is not directly observable. For instance, we can observe a person engage in different behaviors under the same conditions, but the process of action selection is latent. Even if we had the whole wiring diagram

of the person’s brain and spiking patterns under a variety of situations, it still would not be obvious where, how and when action selection takes place (Brown, 2014). So, how to go about it?

Concentrating on behavioral methodology, the study of human cognition can be considered as an instance of a black box problem (Ljung, 1999, 2010). Working on a black box problem, a researcher only has direct information about the system’s input and output, but lacks direct access to internal processes transforming inputs into outputs. The original problem can then be formulated as one of understanding internal processes from known combinations of inputs and outputs, or in psychological terms, uncovering the human cognitive architecture from information about overt behavior taking place in some environment.

A powerful approach to providing an approximate solution to the black box problem is construction of competing mathematical models that can be tested against data using statistical methods (Ljung, 1999, 2010; Casella and Berger, 2002; Gelman et al., 2013a). Each model would represent a mechanistic account of how inputs are transformed into outputs. A model that best balances parsimony, fit to known empirical regularities and interpretability can be taken as a formal instantiation of the best theoretical principles characterizing internal processes and provide further guidance in experimental work. This approach fits well with psychological research because it is often easy to come up with several, categorically different conjectures about underlying cognitive processes (Vickers, 1979; Townsend and Ashby, 1983) and there is plenty of data to constrain them (Luce, 1986).

The primary benefit of developing mathematical models rather than verbal descriptions is that they force coherency on theoretical assumptions, enable deriving exact conclusions from assumptions and improve falsifiability of proposed principles through quantitative statistical testing (Coombs, 1983; Busemeyer and Diederich, 2010; Lewandowsky and Farrell, 2010; Lee and Wagenmakers, 2014). For instance, a model of memory processes that can predict a power relation between time and memory retention is easier to falsify than a verbal account that only makes ordinal predictions (Cavagnaro et al., 2013). Ultimately, the promise of mathematical models is to bring psychology closer to sensible theoretical principles that characterize elementary processes and how they evolve over time, under different experimental conditions for different individuals.

Before presenting an example of the formal modeling framework, I will describe notation used for the rest of the thesis. I will use lower-case non-bolded symbols for scalars, lower-case bolded symbols for column vectors, and upper-case bolded symbols for matrices. I will discriminate between random variables and their realized values by qualifying my statements when necessary, but otherwise their case will reflect their mathematical nature (i.e. scalar or vector), as for other symbols. Observables, response variables or covariates, will be represented with Roman letters while parameters will be represented with Greek letters.

A formal relationship between response variables $\mathbf{y} \in \mathcal{Y} \subseteq \mathbb{R}^k$, covariates $\mathbf{x} \in \mathcal{X} \subseteq \mathbb{R}^l$ and parameters $\theta \in \Theta \subseteq \mathbb{R}^m$, will be stated as a parametric statistical model $M = \{F(\mathbf{y} \mid \mathbf{x}, \theta)\}$, where F is a probability distribu-

tion function. In words, a parametric model is a collection of probability distributions indexed by a finite-dimensional parameter that provides a full description of variability in \mathbf{y} . If a probability density (mass) function f exists, then a model expression can be stated in terms of f . Also, covariates may not always be present, so \mathbf{x} may be dropped. Lastly, a short-hand notation $\mathbf{y} \sim F(\mathbf{y} \mid \boldsymbol{\theta})$ may be used, where \sim stands for “distributed as”.

As an example, consider a random vector $\mathbf{y} = (y_1, y_2, \dots, y_n)^T$, where T is the transpose operator. If we assume that $\mathbf{y} \in \mathbb{R}^n$ varies according to a multivariate normal probability density f , parameterized with a vector $\boldsymbol{\mu} \in \mathbb{R}^n$ and a positive definite matrix $\boldsymbol{\Sigma} \in \mathbb{R}^{n \times n}$, then the parametric model of \mathbf{y} is a collection of probability densities

$$M = \left\{ f(\mathbf{y} \mid \boldsymbol{\mu}, \sigma^2) = \det(\sqrt{2\pi}\boldsymbol{\Sigma})^{-1} \exp \left(-\frac{1}{2}(\mathbf{y} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}(\mathbf{y} - \boldsymbol{\mu}) \right) \right\}, \quad (1.1)$$

where $\det(\cdot)$ is a determinant operator.

With notation fixed, I will briefly present the formalism of the standard signal detection theory as an example of approach taken to testing dependencies among parameters during simple choice behavior (Macmillan and Creelman, 2004). The original phenomenon motivating development of signal detection models was the human ability to make decisions with distorted sensory stimuli. This phenomenon can be studied experimentally using a computerized task. For example, during the auditory signal detection task an observer is presented with a sequence of stimuli drawn either from a *Signal* distribution representing a tone combined with white noise, or a *Noise* distribution representing white noise. On each trial, an observer has to respond

“Yes” if he or she decides the stimulus belongs to the *Signal* distribution and “No” otherwise. Under these conditions, a researcher knows to which class each presented stimulus belongs and observes a sequences of responses, typically summarized as hit rate (proportion of true positive responses) and false alarm rate (proportion of false positive responses). The psychological problem is to characterize cognitive processes involved in transforming stimulus information into responses from the generated data.

Signal detection models are statistical models that decompose hit and false alarm rates into cognitively interpretable parameters characterizing stimulus representation and decision process without specifying mechanistic details (Macmillan and Creelman, 2004; Lee and Wagenmakers, 2014). The basic idea is that behavior in a signal detection task can be explained by a theory that proposes the following principles:

1. Presentation of a stimulus causes the observer to form a representation of the stimulus “strength.”
2. An observer is able to establish a decision threshold to which this perceived strength can be compared.
3. If the perceived strength is greater than this threshold then the observer responds “Yes,” otherwise he or she responds “No.”
4. The perceived strength varies from trial to trial such that its statistical structure (mean and variance) is determined by the class of the presented stimulus, where the strengths of Noise stimuli tend to be per-

ceived as being below the threshold and the strengths of Signal stimuli tend to be perceived as being above.

Formalization of these principles can be done using language of probability theory (Casella and Berger, 2002). Let a random variable $c \in \{0, 1\}$ represent a stimulus randomly drawn either from the *Noise* class ($c = 0$) or the *Signal* class ($c = 1$). The probability mass function of c is determined by the experimenter. In response to the sampled stimulus, an internal representation of strength value, represented by a random variable $x \in \mathbb{R}$, is generated. We commonly assume that x conditioned on the stimulus class is normally distributed with class-specific mean $\mu_i \in \mathbb{R}$ and variance $\sigma_i^2 > 0$ parameters, so that

$$\begin{aligned} x \mid c = 1 &\sim \mathcal{N}(\mu_s, \sigma_s^2) \text{ and} \\ x \mid c = 0 &\sim \mathcal{N}(\mu_n, \sigma_n^2), \end{aligned} \tag{1.2}$$

where \mid stands for “conditioned on”. Finally, let $\tau \in \mathbb{R}$ represent the position of a threshold on a strength dimension. Then a random variable

$$r = \begin{cases} 1 & \text{if } x \geq \tau \\ 0 & \text{if } x < \tau \end{cases} \tag{1.3}$$

represents a “Yes” ($r = 1$) or a “No” ($r = 0$) response.

As written, this version of the model has five free parameters $(\mu_n, \mu_s, \sigma_n, \sigma_s, \tau)$, but because normal distributions form a location-scale family, where the mean and standard deviation are the location and scale parameters, respectively, the means and standard deviations are identifiable only relative to

each other. Therefore, we set $\mu_n = 0$ and $\sigma_n = 1$ without loss of generality. The other three parameters remain free, and quantify representational and decision features of signal detection.

In actual analysis of behavioral data, it is common to use the derived parameter $d' = (\mu_s - \mu_n) / \sqrt{(\sigma_s^2 + \sigma_n^2)/2} = \mu_s / \sqrt{\sigma_s^2/2}$, which is psychologically interpreted as quantifying how well the encoding process separates *Noise* from *Signal* stimuli (Macmillan and Creelman, 2004). The threshold τ represents the decision criterion and the ratio of variances $\sigma_n^2 / \sigma_s^2 = 1 / \sigma_s^2$ between *Noise* and *Signal* distributions characterizes relative noise in representations. These three parameters drive the predicted pattern of responses.

To begin testing a specified cognitive model one needs to derive, or be able to simulate, predictions for behavioral data. In the signal detection case, the random mechanism specified above, and an additional assumption that all responses during an experiment are mutually independent, allow for an analytical derivation of a probability mass function parameterized by the three free parameters. The derived model for the number of hits $y_h \in \{0, 1, \dots, N_s\}$ and the number of false alarms $y_{fa} \in \{0, 1, \dots, N_n\}$ for a given participant presented with N_s *Signal* stimuli and N_n *Noise* stimuli is a product of two binomial probability mass functions. The joint probability of y_h and y_n , conditioned on parameters, is therefore

$$f(y_h, y_{fa} \mid p_h, p_{fa}, N_s, N_n) = \binom{N_s}{y_h} p_h^{y_h} (1 - p_h)^{N_s - y_h} \binom{N_n}{y_{fa}} p_{fa}^{y_{fa}} (1 - p_{fa})^{N_n - y_{fa}}, \quad (1.4)$$

with the probabilities p_h and p_{fa} being the probabilities that $r = 1$ condi-

tioned on the stimulus class. The hit and false alarm probabilities are in turn functions of the representational and decisional parameters μ_s , τ and σ_s , such that

$$p_h = P\{r = 1 \mid c = 1\} = \Phi\left(\frac{\mu_s - \tau}{\sigma_s}\right) \text{ and} \\ p_{fa} = P\{r = 1 \mid c = 0\} = \Phi(-\tau), \quad (1.5)$$

where $\Phi(\cdot)$ is the standard normal distribution function and $P\{\cdot \mid \cdot\}$ is a conditional probability.

When a statistical model relating cognitive parameters to data is specified for a given theory, one can use a rich body of statistical methods to explore its predictions via simulations, to estimate its free parameters from data, and to compare and select from competing models of cognitive processing involved in a given task (Berger, 1997; Casella and Berger, 2002; Gelman et al., 2013a). Connecting the signal detection model to performance data enables a researcher to test underlying principles and make a range of inferences about cognitive processes from behavior. For instance, it is possible to test goodness of fit of a model with $\sigma_s^2 = 1$ against a model where σ_s^2 is a free parameter to determine the noise properties of the stimulus representation. However, the range of possible inferences is limited, so the specified model cannot say anything about (for example) formation and adaptation of the representations to dynamic stimulus stream or motor processing. It would require a new model to address additional questions about processes involved in a signal detection task.

The example of a standard signal detection model shows how experimen-

tal variables and behavioral data can be tied together in a formal model. With a formal model, theory of signal detection can be better tested and behavioral data can be interpreted in a psychologically interesting way. I intend to carry over these analytical advantages to my proposed studies of simple decision making. In the next section, I turn back to the problem of simple decision making and set the stage for later model development by elaborating on the structure of experimental tasks used to obtain behavioral data.

1.2 Data-Generating Experimental Tasks

To obtain input-output data informative of processes underlying simple choice behavior, we need a laboratory task that recruits processes of interest. I will consider simple choice behaviors to be a class of behaviors that rely on a single deliberation stage to form a discrete decision about an objectively determinable state of a noisy stimulus pattern during at most a two second window, starting with stimulus presentation and ending with movement registration. A large collection of tasks can evoke simple choice behaviors and can be classified in terms of five dimensions: the sensory modality, the number of choices, the number of stimulus attributes relevant to a decision, whether a participant or an experimenter controls initiation of responses and response modality. For this thesis, I will concentrate on participant-controlled two-choice tasks, where finger movements are driven by a single attribute of a visual stimulus. For short, I will call them two-choice tasks.

Two-choice tasks have been the predominant kind for obtaining empirical regularities and developing theoretical models of simple decision making (Laming, 1968; Vickers, 1979; Luce, 1986; Ratcliff and Smith, 2004; Ratcliff and McKoon, 2008; Wagenmakers, 2009). An example of a two-choice task is a numerosity task (Ratcliff et al., 2012). On a given trial, two clouds of points, consisting anywhere from a few dots to a hundred, are presented simultaneously on opposite sides of a computer screen. The relative number of points determines uncertainty about which cloud is more numerous. An observer has to make a decision, as accurately and as fast as she can, according to a rule that maps stimulus identity onto movements, such as “press the button on a keyboard corresponding to the cloud with the larger number of points”. Accuracy and response time are measured using a computer.

Across trials of a two-choice task, several experimental variables are usually manipulated that to produce recurrent empirical regularities in responses and response times (Luce, 1986). Continuing with the numerosity task, a researcher could vary relative number of points, instructions emphasizing either speed or accuracy, probability of the “large” cloud appearing on the left side and probability of transitions between “large” and “small”. All of these manipulations tend to affect both accuracy and speed of responses, but in different ways, providing non-trivial constraints to test proposed models (Vickers, 1979; Townsend and Ashby, 1983).

The combination of experimental variables and response variables generated by two-choice tasks make up the input-output data that can be used to learn about cognitive processing. In the next section, I will present a bench-

mark dataset generated by a brightness discrimination task that demonstrates the kind of patterns models need to satisfy.

1.3 Benchmark Behavioral Data

Two-choice tasks are simple, but under different experimental manipulations they generate many informative patterns in behavioral data. I will use the Ratcliff and Rouder (1998) brightness discrimination data as a source of behavioral data to test for dependencies among parameters. This dataset has become somewhat of a benchmark for testing new models because it has a large number of observations, and demonstrates many of the recurrent accuracy and response time patterns (Vandekerckhove and Tuerlinckx, 2008; Brown and Heathcote, 2008; Vandekerckhove et al., 2011; Verdonck and Tuerlinckx, 2014).

Ratcliff and Rouder (1998) ran a brightness discrimination task with three participants. Each participant had one 35 min practice session and ten 35 min experimental sessions. Each session consisted of eight blocks of 102 trials, for a total of 8160 trials.

Each trial began with participants seeing a gray screen for 500 ms. Then, an array of white and black pixels appeared at the center. Participants had to decide whether the array came from a “high” or “low” brightness distribution. They responded as soon as they could by pressing one of two buttons with the corresponding finger. Upon a response, the stimulus would change back to the gray screen for a duration of 300 ms. During the final

300 ms, a feedback message was displayed indicating whether the decision was correct or not. Both responses and response times were recorded upon completion of a trial.

During each session, two experimental variables were manipulated. First, the ratio of white to black pixels, representing stimulus difficulty, was manipulated across trials. On each trial, a ratio was drawn from one of two symmetric, unimodal probability mass functions spanning over 33 levels of brightness and overlapping to ensure imperfect accuracy. Second, instructions changed every 204 trials to direct participants to either maximize speed or maximize accuracy.

Overall, with a simplifying assumption that practice effects are negligible after the training session, the experiment generated a data matrix \mathbf{D} with 24,480 rows observations, and 5 columns of variables. Among the five variables, response times and choices will be treated as response variables, and participant, brightness level and instructions will be treated as covariates. In the rest of this section, I will examine patterns in response variables under different values of covariates.

One important feature of behavioral data is large variation in response times within as well as between participants. Figure 1.1 shows estimated error and correct response time probability densities for two participants (“nh”, “jf”) placed on the negative and positive half of the real number line, respectively. Response times come from trials when each participant was presented with the same, very difficult brightness level (0.469 white to black pixels), and instructed to be accurate. Response times from “high”

and “low” stimuli were symmetric, and thus plotted functions combine data for both stimuli. Resulting probability densities show a large spread and a positive skew for both participants. The two participants, however, show considerable quantitative differences. For example, Participant 1 (black line) has a mean correct response time of 743 ms relative to 815 ms for Participant 2 (blue line). Other features like skew, standard deviation and kurtosis also show strong individual differences.

Another central feature of behavioral data is its asymmetry between correct and error responses. Such an asymmetry may be seen in Figure 1.1, where error means are slower than correct means, which is typical for difficult stimulus conditions (Swensson, 1972). However, a more complete understanding of the relation between correct and error response times arises from examining performance across several difficulty and instruction combinations. One effective graphical summary that can reveal relations between correct and error response times is called a latency-probability graph (LPG). The LPG is a function that maps proportions of correct or error responses into mean response times, thus tracking stimulus difficulty conditions (Ratcliff and Rouder, 1998).

Figure 1.2 shows a plot of a LPG for one participant presented with stimuli of different brightness, either under speed or accuracy instructions. When accuracy was emphasized, the participant showed faster errors for easy stimuli (really bright or dark) and slow errors for difficult stimuli relative to correct responses. When instructions shift to *speed* the asymmetry is strongly reduced or vanishes (Vickers, 1979).

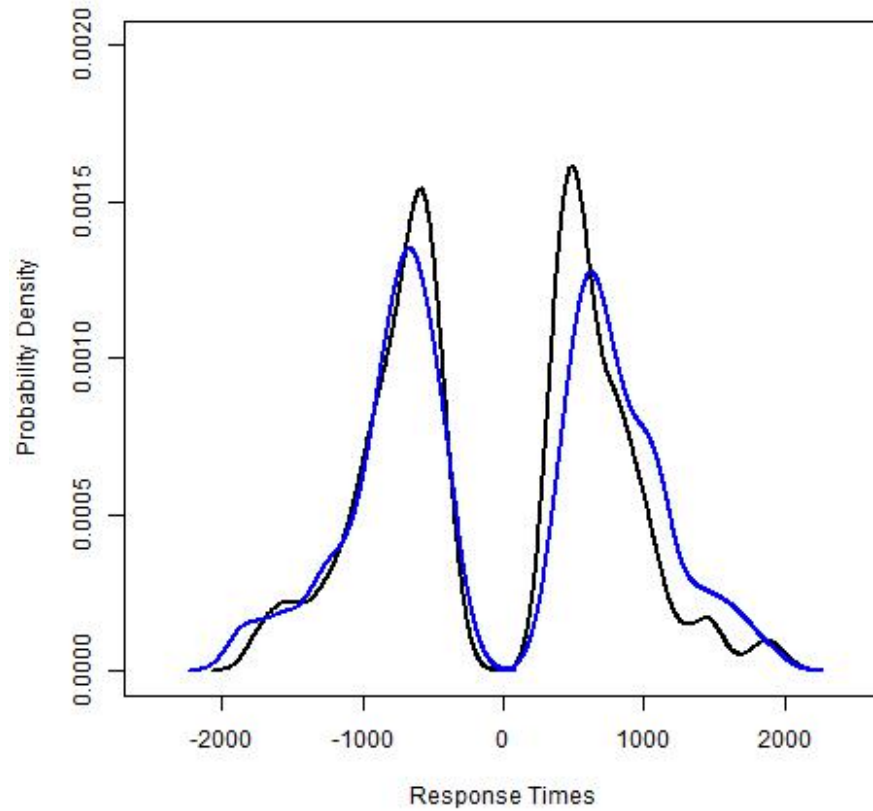


Figure 1.1: Estimated correct and error probability densities of response times. Positive response times correspond to correct choices and negative response times to incorrect choices. The black line is Participant “nh” and the blue line is Participant “jf”.

Another pattern that is ubiquitous in behavioral data is the speed-accuracy trade off obtained by manipulating instructions (Vickers, 1979; Luce, 1986; Heitz, 2014). Instructing a participant to emphasize accuracy instead of speed results in shifting and rescaling both correct and error densities towards slower response times. The effect is obvious in Figure 1.2 with the LPG shifting upwards and to the right when instructions are changed from “accuracy” to “speed”.

In contrast to the speed-accuracy trade-off under instructions, concentrating on either the speed or accuracy condition, the relation between response times and accuracy is different. Figure 1.2 shows that the higher accuracy the is, the faster the participant responds. Intuitively, when the stimulus gets easier to discriminate, the participant is faster and more accurate.

The last two aspects of the benchmark data are shown in Figure 1.3. The plot is based on response times collected in two consecutive blocks under speed instruction. Similar data features appear in other blocks. The first feature is the presence of extreme response times. For example, the Trial Series subplot shows that Response 9 took 538 ms, and Response 21 lasted 196 ms. Relative to the rest of the responses, these have extreme duration. Presence of extreme observations may indicate genuine response times from tails of the distribution carrying information about processing of interest, but it is not completely clear (Craigmile et al., 2010).

The second feature is portrayed in the Auto-correlation subplot of Figure 1.3. The dotted blue line represents a critical value for an alpha-level of 0.05 hypothesis test of zero auto-correlation. Data from the selected block is

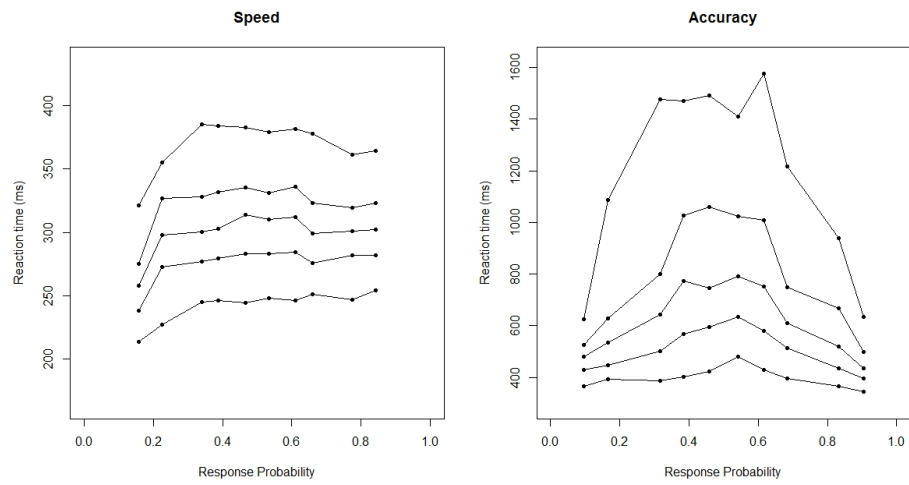


Figure 1.2: Latency-probability plot. Below 0.5 are error rates and above 0.5 are correct rates. Points represent a summary of performance with lines drawn to capture a trend. The lower line is due to speed emphasis and the upper line is due to accuracy emphasis.

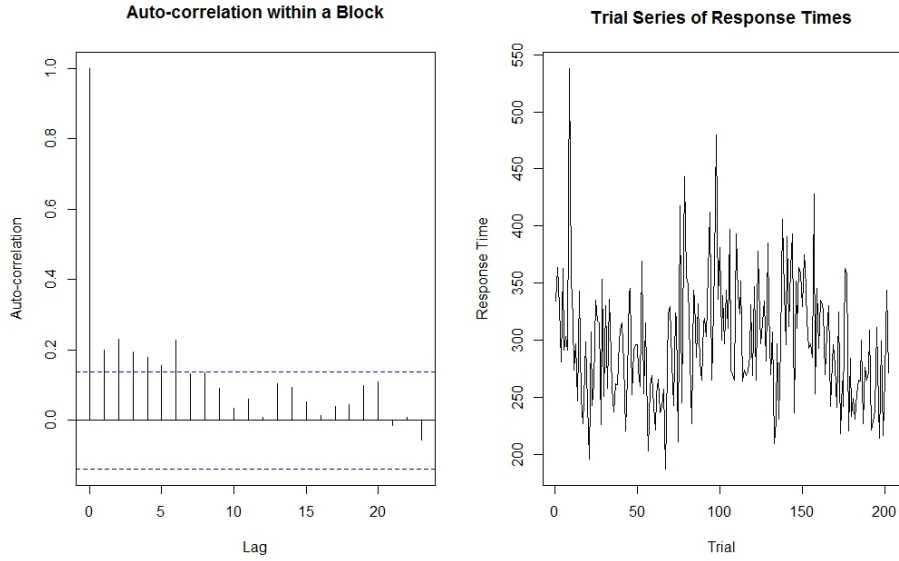


Figure 1.3: Auto-correlation function and trial series of response times. Values above the dotted line represent significant auto-correlation.

consistent with auto-correlation between response times stretching 6 trials.

The complete pattern of data features including speed-accuracy tradeoff, amount of variation in response times, skew of the conditional densities, relation between correct and error response times, relation between response time and accuracy, extreme response times and auto-correlation in response times provides strong constraints for developing and testing models of decisions. In the next section, I will discuss a prominent theoretical framework for understanding simple choice behavior and how to implement it in testable models.

1.4 Theoretical Data Models

A common approach to conceptualizing cognitive processing underlying two-choice tasks is to decompose them into the three processes of (1) encoding of stimulus features, (2) response selection based on a relevant stimulus feature and (3) motor execution. A variable cognitive process model for these tasks must be able to predict patterns in performance data. A delicate balance has to be struck between what processing details to include and what details to exclude from models because it will determine their generalizability across two-choice tasks. For instance, including details about how lexical information is accessed in memory to direct response selection will make a model inapplicable to brightness discrimination. The dominant approach to this balancing problem is the sequential sampling framework.

A fundamental theoretical assumption of the sequential sampling framework is that the same general decision process operates across all two-choice tasks. Then, a generalizable model should provide a detailed account of the decision process while making simplifying assumptions about the stimulus encoding and motor processing, like assuming a constant non-decision time. Over the last approximately 50 years, many models of this type have been developed (Stone, 1960; Ratcliff, 1978; Vickers, 1979; Ratcliff and Tuerlinckx, 2002; Smith, 1995; Usher and McClelland, 2001; Brown and Heathcote, 2008). Some of the most successful ones can explain many patterns in the response time and accuracy data, under a variety of experimental conditions and tasks (Ratcliff and Smith, 2004; Wagenmakers, 2009).

The underlying principle of sequential sampling models is that when an observer is tasked with selecting a correct response in limited time from a noisy stimulus, he or she accumulates samples of information from the relevant stimulus attribute representation until one of two decision thresholds is reached (Vickers, 1979; Townsend and Ashby, 1983; Luce, 1986; Bogacz et al., 2006). The structure of thresholded accumulation unites concepts of decision bias, caution and information quality under a single mechanism. Differential levels of initial information produce bias, thresholds control caution and rate of information uptake reflects information quality. Combinations of these variables in thresholded accumulation also naturally predict decisions and response times. The time taken to reach a threshold and the crossed threshold determine decision time and decision. The sequential sampling principle is powerful, but to connect it to behavioral data we need a model.

Given the sequential sampling principle, any concrete model of simple decisions makes detailed assumptions about the nature of sampled information, temporal structure of sampling, stopping rule and sources of randomness (Vickers, 1979; Ratcliff and Smith, 2004; Bogacz et al., 2006; Teodorescu and Usher, 2013). For example, sampled information can consist of discrete or continuous packets. The size of packets may be fixed or variable. The accumulation of information packets may happen at fixed or variable discrete time points, or continuously. The stopping rule may be absolute, in which case accumulation stops when evidence in favor of one of the choices reaches a threshold, regardless of the amount of evidence favoring the alternative. In contrast, a relative rule requires that evidence favoring one of the

choices has to exceed the other by some amount. Finally, models vary in whether starting points, rate of information uptake or thresholds vary across trials. Combinations of modeling assumptions generate the range of possible sequential sampling models.

Examples of sequential sampling models include a random walk model that stops accumulation of discrete packets of information over discrete time steps according to a relative rule (Laming, 1968), a Poisson counter model that describes accumulation of discrete, fixed information packets arriving at random time points until evidence for either of the choices reaches a threshold (Townsend and Ashby, 1983), and a Wiener diffusion model that describes accumulation of continuous packets in continuous time until a relative threshold is hit (Ratcliff, 1978).

In addition to assumptions about the accumulation process, models also add assumptions about variability in their parameters across trials. For example, an early random walk model assumed variability in starting points (Laming, 1968), a version of a Poisson counter model reported in Ratcliff and Smith (2004) added variability in the rate of information uptake and thresholds, and a popular version of a Wiener diffusion model incorporates variability in starting points, rate of information uptake and non-decision time.

The sequential sampling framework has been fruitful in providing models that could be used to investigate correlations between parameters. The assumptions of a particular model lead to a precise prediction of how response times and responses jointly vary (Ratcliff and Smith, 2004). The joint density

can be used as a theoretically motivated data model in statistical analysis of behavioral data. In the next section, I will characterize a Bayesian statistical framework that allows a principled and intuitive way for extracting information about the supposed latent decision process from behavioral data.

1.5 Bayesian Statistical Framework

Once we form a model of behavioral data, how should we use the data? In the following subsections, I will present a Bayesian hierarchical framework and a fitting method based on Markov chains that can provide a formal approach to examining generalized models of simple choice behavior (Berger, 1997; Casella and Berger, 2002; Gelman et al., 2013a). In the Bayesian framework a model of data is combined with a model of unknown parameters to provide a combined description of uncertainty. While theoretically everything tends to work out beautifully, in practice Bayesian models with complex data distributions are usually analytically intractable and require an iterative algorithm, often called a sampler, to learn about parameters and the model(s) from data.

Recent developments in computing power and software made the Bayesian approach ripe for scientific use (Kruschke, 2010), and it is becoming popular in psychology, too (Edwards et al., 1963; Myung and Pitt, 1997; Peruggia et al., 2002; Rouder et al., 2005; Rouder and Lu, 2005; Myung et al., 2008; Craigmile et al., 2010; Vandekerckhove et al., 2011; Lee and Wagenmakers, 2014). The reasons for its appeal include quantifying of uncertainty

with probability, making statistical inferences as probability statements and gaining the ability to fit realistic models. On these grounds, I will adopt Bayesian approach to a data analysis study specified below, but before that I will present the fundamental ideas of hierarchical modeling and Markov Chain Monte Carlo.

1.5.1 Bayesian Hierarchical Models

In the Bayesian framework model parameters are treated as random variables along with data. The full probability model, also called the Bayesian model, is at the center of analysis because all inference problems can be defined in terms of it (Gelman et al., 2013a). A Bayesian model requires specifying the probability densities of both data and parameters. The probability densities are assumed to be conditionally dependent in a way that the full joint density can be factored into a probability density of data given parameters, called a sampling density, and the probability density of parameters, called a prior density. Given a data vector $\mathbf{x} \in \mathbb{R}^p$ and its parameter vector $\boldsymbol{\theta} \in \Theta \subseteq \mathbb{R}^m$, a Bayesian model is written as

$$f(\mathbf{x}, \boldsymbol{\theta}) = f(\mathbf{x} \mid \boldsymbol{\theta})\pi(\boldsymbol{\theta}) \quad (1.6)$$

where $f(\mathbf{x} \mid \boldsymbol{\theta})$ is a sampling density and $\pi(\boldsymbol{\theta})$ is a prior density for $\boldsymbol{\theta}$. The ability to factor a joint density of \mathbf{x} and $\boldsymbol{\theta}$ into a sampling and a prior density is useful for constructing new Bayesian models because it allows breaking the problem into manageable pieces, often consisting of common univariate distributions.

Statistical inference proceeds by conditioning unknown quantities of interest on the observed data \mathbf{x} . For example, estimation of parameters $\boldsymbol{\theta}$ amounts to calculating a posterior density $\pi(\boldsymbol{\theta} \mid \mathbf{x})$, which contains all the information about the parameters. Using Bayes’ theorem,

$$\pi(\boldsymbol{\theta} \mid \mathbf{x}) = \frac{f(\mathbf{x} \mid \boldsymbol{\theta})\pi(\boldsymbol{\theta})}{\int_{\Theta} \dots \int f(\mathbf{x} \mid \boldsymbol{\theta})\pi(\boldsymbol{\theta})d\boldsymbol{\theta}} \quad (1.7)$$

(Berger, 1997; Casella and Berger, 2002; Bernardo, 2011; Gelman et al., 2013a). We can then extract more specific probability statements about $\boldsymbol{\theta}$ from the posterior density such as $P(\boldsymbol{\theta} \in A \mid \mathbf{x})$, where $A \subseteq \Theta \subset \mathbb{R}^m$.

The expression in Equation 1.6 suggests a situation where a set of observations all come from the same source and can be used to infer something about common parameters. However, the structure of behavioral data is rather different. A typical experiment generates multiple observations for each of several participants engaged in the same task (Ratcliff and McKoon, 2008; Wagenmakers, 2009). When specifying a statistical model, the experimenter must choose whether to treat individual observations as coming from different distributions, a “no pooling” model, or the same distribution, a “complete pooling” model (Rouder et al., 2005; Rouder and Lu, 2005; Rouder et al., 2014).

Not pooling the data presupposes no relation between the individuals while completely pooling the data ignores individual differences. However, for human behavioral data, it is more plausible to assume a middle ground. Real data varies among participants, but due to similarity in cognitive abilities and experimental treatment, participants share regularities. A hierar-

chical statistical model acknowledges this mixed structure by assuming that individual parameters are a random sample from a population with data nested within each individual, and thus introduces “semi-pooling” of information by fitting all the data simultaneously (Gelman et al., 2013a). Hence, I propose treating sequential sampling models under consideration in a hierarchical fashion (Berger, 1997; Gelman et al., 2013a; Rouder et al., 2005; Craigmile et al., 2010).

There is no rigorous definition of a hierarchical model, but it is often considered that a model with at least two levels of parameters qualifies as hierarchical (Berger, 1997; Gelman et al., 2013a). In the simplest case, the parameter space can be factored into first-level parameters representing individuals and second-level parameters representing a population from which individuals’ parameters are sampled. With data $\mathbf{x} \in \mathbb{R}^p$ conditionally dependent on the first-level parameters $\boldsymbol{\theta} \in \Theta \subseteq \mathbb{R}^m$, and first-level parameters conditionally dependent on second-level parameters $\boldsymbol{\psi} \in \Psi \subseteq \mathbb{R}^n$, a two-level hierarchical model can be stated as

$$f(\mathbf{x}, \boldsymbol{\theta}, \boldsymbol{\psi}) = f(\mathbf{x} \mid \boldsymbol{\theta})f(\boldsymbol{\theta} \mid \boldsymbol{\psi})f(\boldsymbol{\psi}). \quad (1.8)$$

In words, a joint density over data and all parameters can be factored into a sampling density $f(\mathbf{x} \mid \boldsymbol{\theta})$, a first-level prior density of parameters $\pi(\boldsymbol{\theta} \mid \boldsymbol{\psi})$ and a second-level prior density of parameters $f(\boldsymbol{\psi})$. In the context of sequential sampling models, a density over a collection of responses and response times is a sampling density, a density over individuals correlation parameters between bias and non-decision time, say, is the first-level prior and a density

over the correlation parameter is the second-level prior. Combining all the data together hierarchically and then conditioning on all of it will lead to improved estimation of important cognitively interpretable parameters of each individual because data is shared across individuals. Sharing of information leads to what is sometimes called a shrinkage effect (Gelman et al., 2013a).

While hierarchical Bayesian models have many advantages, one necessary assumption taken when using hierarchical models, which I will adopt, is exchangeability of observations. A sequence of random variables $\{x_1, x_2, \dots, x_n\}$ is exchangeable when

$$f(x_1, x_2, \dots, x_n) = f(x_{z(1)}, x_{z(2)}, \dots, x_{z(n)}), \quad (1.9)$$

for all permutations z over $\{1, 2, \dots, n\}$ (Gelman et al., 2013a). In other words, order indices of random variables carry no information about them. The consequence of exchangeability is that, conditional on first level parameters $\boldsymbol{\theta}$, observations are independent and a prior for $\boldsymbol{\theta}$ exists, formally justifying Bayesian analysis. Exchangeability is a simplifying assumption because actual response times have an autocorrelation structure that is altered by permutations, as Figure 1.3 shows. However, numerous successful fits of sequential sampling models suggest that exchangeability is a reasonable simplification (Usher and McClelland, 2001; Ratcliff and Smith, 2004; Ratcliff and McKoon, 2008).

To conclude, the joint density over observations and all parameters is the starting point of Bayesian analysis. I will use a hierarchical modeling approach to capture the patterns of similarities and differences between hu-

mans performing simple decision tasks, while simplifying the models with assumption of exchangeability. However, the models I will define are still too complex for analytical treatment, which brings us to the topic of how complicating the situation with Markov chains actually enables practical application of Bayesian inference.

1.5.2 MCMC Background

After a Bayesian model has been specified, inference about parameters, out-of-sample observations or the model can proceed by applying the calculus of probabilities to obtain corresponding conditional densities. With continuous parameters, which is the typical case, calculating conditional densities requires integration, which for complex models is intractable. Application of Bayes' theorem to obtain a posterior density over unknown quantities is complicated by the integral in the denominator of Equation 1.7. In particular, the inability to integrate over parameters relative to their density plagues the best sequential sampling models, so the posterior density cannot be calculated exactly (Rouder et al., 2005; Vandekerckhove et al., 2011).

Due to the high-dimensionality of Bayesian models, numerical integration techniques can fall apart, leaving Monte Carlo methods as the only viable option. Monte Carlo methods rely on simulating pseudo-random draws from a distribution of interest to solve inference problems. For data models with known mathematical form, approximate inferences can be obtained through a special class of Monte Carlo algorithms, sometimes called samplers, that

use Markov chains (MCMC) to sample from posterior densities (Robert and Casella, 2004; Gamerman and Lopes, 2006; Givens and Hoeting, 2012; Gelman et al., 2013a). Given a posterior sample, accurate conclusions can be drawn, even for complicated models (Peruggia et al., 2002; Craigmile et al., 2010; Vandekerckhove et al., 2011). In the rest of the section, I will present a few definitions from Markov chain theory as background for a general class of MCMC algorithms built on that theory.

MCMC employs discrete-time, continuous state space Markov chains to sample from arbitrary distributions for which a probability density can be specified (Karlin and Taylor, 1975, 1981; Ross, 2014). Let $\mathbf{X} = \{\mathbf{x}_n : n \geq 0\}$ be a p -dimensional, discrete-time stochastic process with state space $\mathcal{S} \subseteq \mathbb{R}^p$. A first-order Markov chain is a stochastic process with a conditional distribution satisfying

$$f(\mathbf{x}_n \mid \mathbf{x}_{n-1}, \mathbf{x}_{n-2}, \dots, \mathbf{x}_0) = f(\mathbf{x}_n \mid \mathbf{x}_{n-1}), \quad (1.10)$$

for any $n \geq 1$. In words, the distribution of a current state depends only on the previous state, giving Markov chains a memoryless property.

A Markov chain \mathbf{X} can be thought of as a model of stochastic dynamics of a multi-dimensional system over discrete time. Description of dynamics requires an initial position in the state space and a transition rule that guides the system from state to state. Given that \mathcal{S} is continuous and \mathbf{X} is stochastic, a description of the process requires a marginal density of \mathbf{x}_0 , called an initial density, and a conditional density of \mathbf{x}_n given \mathbf{x}_{n-1} , called a transition density. Once densities describing dynamics are defined, evolution of the

process can be analyzed analytically or through simulation.

It can be shown that some Markov chains, if given enough time, will go from an arbitrary initial density to an equilibrium density. Equilibrium behavior is a property of a transition density. For a transition density $f(\mathbf{x}_n | \mathbf{x}_{n-1})$ and $A \subseteq S$, a transition kernel is defined as

$$P\{\mathbf{x}_n \in A | \mathbf{x}_{n-1}\} = \int_A f(\mathbf{x}_n | \mathbf{x}_{n-1}) d\mathbf{x}_n. \quad (1.11)$$

Starting from any point $\mathbf{x}_0 \in S$, a Markov chain has an equilibrium density if an n th-step transition kernel satisfies

$$\lim_{n \rightarrow \infty} P(\mathbf{x}_n \in A | \mathbf{x}_0) = \pi(A), \quad (1.12)$$

where $\pi(\cdot)$ is called an equilibrium density function.

Once a chain attains equilibrium behavior, whatever the initial position was and no matter how many transitions occur after, the chain remains in equilibrium. The sequence of states $\mathbf{x}_n, \mathbf{x}_{n+1}, \dots$ will be identically distributed according to the equilibrium density function. The equilibrium behavior provides a natural point of contact with the problem of Bayesian inference for complex models. If the equilibrium density is equivalent to the posterior density, then a sample of states of a Markov chain in equilibrium is a sample from the posterior.

It follows then that Bayesian analysis can proceed if a Markov chain can be simulated for which there exists a unique equilibrium density equivalent to the posterior density that can be reached from an arbitrary starting point. The limit-based definition of the equilibrium density implies that it is not

guaranteed to exist nor to be unique. Four conditions need to be satisfied to obtain the desired subclass of Markov chains.

First, a Markov chain has to be homogeneous. Homogeneity is defined by a time-invariant transition density, such that

$$f(\mathbf{x}_n | \mathbf{x}_{n-1}) = f(\mathbf{x}_1 | \mathbf{x}_0), \quad (1.13)$$

for any $n \geq 1$. A homogeneous chain has fixed dynamics across time, but marginal densities $f(\mathbf{x}_n)$ are not guaranteed to be equal across time nor is there a guarantee that marginal densities converge to an equilibrium density. Homogeneity is a flexibility condition allowing for different initial densities and the possibility of convergence to an equilibrium density.

Homogeneous Markov chains have to be constrained further to ensure convergence to a unique equilibrium density. The second requirement is that a Markov chain is ϕ -irreducible. Let the time of first re-visit to A from \mathbf{x}_0 be a random variable $t_A = \inf_{n \geq 1} \{\mathbf{x}_n \in A\}$, where set $A \subseteq S$ and starting point $\mathbf{x}_0 \in A$. A Markov chain is ϕ -irreducible if there exists a marginal density $f(t_A)$, such that for all $\mathbf{x}_0 \in A$ and all $A \subseteq S$,

$$P\{t_A < \infty | \mathbf{x}_0\} > 0. \quad (1.14)$$

ϕ -irreducibility ensures that, no matter where the chain starts, it is possible to return there in a finite amount of time. In other words, there are no regions where a chain gets stuck, which would create a possibility for multiple equilibrium densities.

In addition to moving around the state space without getting stuck in some region, a chain should also have no cyclical behavior. The third re-

quired property of chains with unique equilibrium density is aperiodicity. An aperiodic Markov chain can move from any $A \subseteq S$ back to A in one transition, that is for $n \geq 1$ and $\mathbf{x}_{n-1} \in A$

$$P\{\mathbf{x}_n \in A \mid \mathbf{x}_{n-1}\} > 0. \quad (1.15)$$

As a result, an aperiodic chain can remain in region A for multiple steps.

The first three properties ensure that a Markov chain converges to a unique equilibrium density from most initial positions. To ensure that a process converges from all initial positions, it must be Harris recurrent. For $A \subseteq S$ with $\pi(A) > 0$ and all $\mathbf{x} \in A$, a ϕ -irreducible Markov process with an equilibrium density $\pi(\cdot)$ is Harris recurrent if

$$P\{t_A < \infty \mid \mathbf{x}_0\} = 1. \quad (1.16)$$

Harris recurrence guarantees that all realized chains come back to the region where they start in finite time.

Recall that we started with a problem of statistical inference with complex Bayesian models that can be solved with Monte Carlo methods. Conditions of homogeneity, ϕ -irreducibility, aperiodicity and Harris-recurrence are sufficient to guarantee that a proposed MCMC algorithm can solve intractable integrals in Bayesian inference. The next issue we consider is how to find such algorithms.

1.5.3 Metropolis-Hastings Algorithm

The Markov chain theory on its own provides only conditions that a valid algorithm should satisfy, but gives no clues toward a working algorithm. As

opposed to analysis in probability theory, where having a transition density and figuring out its equilibrium density is the problem, in Bayesian analysis the equilibrium density is known and a transition density needs to be figured out. A very general way to construct transition densities is a class of algorithms called Metropolis-Hastings (Robert and Casella, 2004; Givens and Hoeting, 2012; Gelman et al., 2013a). The algorithm provides a way to define a transition density satisfying the four properties required for a Markov chain to converge to a known equilibrium density from an appropriate initial density.

In the situation where MCMC methods apply, we are in possession of a target density, $\pi(\cdot)$. We need a transition density, $p(\mathbf{y} \mid \mathbf{x})$ (or the associated transition kernel), whose equilibrium density is $\pi(\cdot)$. A sufficient condition for $p(\mathbf{y} \mid \mathbf{x})$ to be a proper transition density is called detailed balance (also called time reversibility),

$$\pi(\mathbf{x})p(\mathbf{y} \mid \mathbf{x}) = \pi(\mathbf{y})p(\mathbf{x} \mid \mathbf{y}). \quad (1.17)$$

Detailed balance says that starting at \mathbf{x} and moving to \mathbf{y} happens as often as starting at \mathbf{y} and moving to \mathbf{x} . In addition, if the transition density $p(\mathbf{y} \mid \mathbf{x})$ has support equal to that of $\pi(\cdot)$, or is contained in it, then it will satisfy properties of homogeneity, ϕ -irreversibility, aperiodicity and Harris recurrence, ensuring convergence from any starting point in the state space to a unique equilibrium density.

Naive or ad-hoc methods for constructing a transition density are likely not to satisfy detailed balance, except for maybe special cases, but their

failure is instructive. When a method fails to satisfy detailed balance it implies that, for some combinations of \mathbf{x} and \mathbf{y} , $\pi(\mathbf{x})p(\mathbf{y} | \mathbf{x}) > \pi(\mathbf{y})p(\mathbf{x} | \mathbf{y})$ or the reverse. Loosely, the chain moves too frequently from \mathbf{x} to \mathbf{y} and rarely from \mathbf{y} to \mathbf{x} . The Metropolis-Hastings algorithm is a general approach to MCMC algorithms that provides a way to correct the imbalance by introducing a possibility that a transition may be rejected.

Suppose the chain at time $n - 1$ is in state $\mathbf{x}_{n-1} = \mathbf{x}$. For a move to take place, a Metropolis-Hastings algorithm requires a proposal density $q(\mathbf{y} | \mathbf{x})$ from which a candidate value can be drawn. A candidate value is then accepted with probability $\alpha(\mathbf{y}, \mathbf{x})$, called the acceptance probability. The transition density in a situation when the chain moves ($\mathbf{x}_n = \mathbf{y}$) is

$$q(\mathbf{y} | \mathbf{x})\alpha(\mathbf{y}, \mathbf{x}). \quad (1.18)$$

In a situation where the chain does not move ($\mathbf{x}_n = \mathbf{x}_{n-1}$), there is a probability mass at \mathbf{x} ,

$$r(\mathbf{x}) = 1 - \int_S q(\mathbf{y} | \mathbf{x})\alpha(\mathbf{y}, \mathbf{x})d\mathbf{y}, \quad (1.19)$$

with the integral accounting for all possible moves. Combining the two pieces, for $d\mathbf{y} \subseteq S$ and $\mathbf{x} \in S$, the transition kernel of a Metropolis-Hastings chain is

$$P\{\mathbf{x}_n \in d\mathbf{y} | \mathbf{x}_{n-1}\} = q(\mathbf{y} | \mathbf{x})\alpha(\mathbf{y}, \mathbf{x})d\mathbf{y} + r(\mathbf{x})\delta_{\mathbf{x}}(d\mathbf{y}), \quad (1.20)$$

where $q(\mathbf{y} | \mathbf{x}) = 0$ when $\mathbf{y} = \mathbf{x}$, and Dirac's $\delta_{\mathbf{x}}(d\mathbf{y}) = 1$ when $\mathbf{x} = \mathbf{y}$ and 0 otherwise.

The final part of the algorithm is the definition of the acceptance probability. From the consideration of the two kinds of inequalities that violate

the detailed balance condition, the acceptance probability for moving from \mathbf{x} to proposed \mathbf{y} is

$$\alpha(\mathbf{y}, \mathbf{x}) = \begin{cases} \min\left(\frac{\pi(\mathbf{x})q(\mathbf{y} | \mathbf{x})}{\pi(\mathbf{y})q(\mathbf{x} | \mathbf{y})}, 1\right) & \text{if } \pi(\mathbf{y})q(\mathbf{x} | \mathbf{y}) > 0 \\ 1 & \text{if } \pi(\mathbf{y})q(\mathbf{x} | \mathbf{y}) = 0 \end{cases}$$

The form of $\alpha(\mathbf{y}, \mathbf{x})$ implies that the target density needs to be known only up to a normalizing constant which cancels out in the fraction. The cancelation of normalizing constants is directly relevant for Bayesian inference because these constants are the intractable integrals. The form also implies that the algorithm always moves to places of higher probability density, but moves to places of lower probability with some probability. Algorithm 1 shows some Metropolis-Hastings pseudocode for sampling N draws from a target density $\pi(\cdot)$.

Algorithm 1 Metropolis-Hastings Pseudocode

Initialize at \mathbf{x}_0 such that $\pi(\mathbf{x}_0) > 0$

for $j = 1, 2, \dots, N$ **do**

 Sample $\mathbf{y} \sim q(\mathbf{y} \mid \mathbf{x}_{j-1})$

 Sample $u \sim \mathcal{U}(0, 1)$

if $\alpha(\mathbf{y}, \mathbf{x}_{j-1}) \geq u$ **then**

 Set $\mathbf{x}_j = \mathbf{y}$

else

 Set $\mathbf{x}_j = \mathbf{x}_{j-1}$

end if

end for

Chapter 2

Questions to be Answered

In the first chapter I raised some questions about dependent adjustment in parameters of cognitive processes underlying simple choice behavior. I presented some behavioral and modeling evidence for it, and reviewed a theoretical and statistical framework within which answers to these questions may be sought. In this chapter, I will elaborate on questions to be addressed by this thesis. There are three overarching questions that I propose to pursue: How to incorporate trial-to-trial dependencies between parameters into models of cognitive processes underlying behavior on two-choice tasks? How do predictions of cognitive models change under assumption of dependence relative to independence of parameters? What information about dependencies between parameters, given some sequential sampling model of the decision process, does the benchmark dataset have?

2.1 Modeling Parameter Dependencies

One of the main questions underlying this thesis is how to formally describe coordination of cognitive processes across trials. However, before stating the question in an answerable form, recall that behavioral and modeling data suggest that across trials parameters that capture underlying processes show both auto-correlation and cross-correlation. As a final target of research we may consider finding a system of equations that can describe such dynamics in a way that reflects a mechanistic theory.

Suppose we are dealing with a sequential sampling model of binary choice that has some initial evidence ξ , rate of evidence accumulation δ , threshold separation α and non-decision time τ . On each trial a participant generates a response time t^{rt} and a response r . Then, change in the task-specific processing from trial n to $n + 1$ may be captured by the following dynamic system:

$$\begin{aligned}\alpha(n + 1) &= f_\alpha(\alpha(n), t^{rt}(n), r(n), \theta_\alpha), \\ \xi(n + 1) &= f_\xi(\xi(n), t^{rt}(n), r(n), \theta_\xi), \\ \delta(n + 1) &= f_\delta(\delta(n), t^{rt}(n), r(n), \theta_\delta), \text{ and} \\ \tau(n + 1) &= f_\tau(\tau(n), t^{rt}(n), r(n), \theta_\tau).\end{aligned}\tag{2.1}$$

The dependence of parameters at trial $n + 1$ on themselves at trial n and other stable processes captured by θ_i can explain first-order auto-correlation, and dependence on previous choice and its speed can explain cross-correlations.

There have been some quantitative proposals that capture auto-correlation

of some parameters across trials, which is a special case of the above system (Laming, 1968; Vickers, 1979; Luce, 1986; Cho et al., 2002; Brown et al., 2008; Gao et al., 2009; Goldfarb et al., 2012). The thrust of this work relies on the rich pattern of sequential effects that provide suggestive constraints on dynamic rules for various parameters. For example, Goldfarb et al. (2012) examined post-error slowing and suggested a dynamic rule for threshold separation that leads to increases in response time after error and decreases after correct responses. Another suggestive sequential effect is the speed up of correct responses on the second repeating stimulus with short response-to-stimulus intervals, which Gao et al. (2009) described with a bias in the starting evidence towards the correct boundary.

However, when it comes to cross-correlations of parameters, the behavioral data is less clear and I found no explicit modeling studies of cross-correlations. Based on these suppositions, I propose to study cross-correlations among parameters from a descriptive rather than a mechanistic perspective. Instead of specifying dynamic rules for all or some of the parameters to impose dependencies, I can replace assumptions of independent variation of the standard models with dependent variation. Dependent variation can be specified as a multivariate distribution that provides means and scales as the standard models, but also adds correlations between parameters. The magnitude and direction of correlations between different pairs of parameters could then suggest a target for a comprehensive mechanistic theory expressed in a system of dynamic equations.

By formalizing parameter relations in a statistical manner, the first ques-

tion can now be formulated more concretely: what multivariate density should model joint variation among parameters across trials? A couple of key properties of standard sequential sampling models provide important constraints on the search space. One feature of all standard models is that parameters fall on different scales. Rates of evidence uptake take values from \mathbb{R} , evidence thresholds are positive, and initial evidence may fall between 0 and the evidence threshold (Ratcliff and Smith, 2004). In addition, probability density functions that we assign to each parameter may all be different to reflect different kinds of marginal variation (Ratcliff, 2013; Jones and Dzhafarov, 2014).

The combination of scales and marginal densities restrictions may be hard to satisfy with common multivariate densities. There is a stock of multivariate densities that could be used to model dependence in certain situations, including the normal, t, gamma, and F (Kotz et al., 2004). However, they are too restricted in their properties to be used for generalizing sequential sampling models. For example, in each of the commonly used densities the support of each dimension is the same and so are the marginal densities. An n -dimensional multivariate normal has support over \mathbb{R}^n and each marginal is a normal density with some mean and variance. These considerations suggest that we need a method for defining new multivariate densities that fit the above restrictions and express alternative correlation structure.

2.2 Predictions of Dependent Parameters

Assuming that we have specified a joint distribution of parameters, we can now examine the consequences of dependent parameters. Assessing predictions made by models is a natural way of understanding what model assumptions imply. Comparison of a model with dependent parameters to a model with independent parameters can reveal how trade-offs between parameters shape the decision process and provide insight into which data patterns indicate dependent parameters. The result of this should be a better understanding of how a decision process with parameters changing across trials in a coordinated manner differs from the usually assumed process with independent parameters.

Sequential sampling models of a simple decision process provide two avenues for understanding implications of correlated parameters. On the one hand, because sequential sampling models are dynamic, we can consider how dependencies alter the accumulation process. Specifically, we can examine average evolution of the accumulation process for a sample of dependent parameters to understand how it changes relative to a model with independent parameters. In addition to theoretical insight, examination of the average evolution may also prove relevant for empirical reasons because dynamics of accumulation can be related to single-cell recording (Ratcliff et al., 2007), functional magnetic resonance imaging (Hare et al., 2011) or electroencephalography (Polanía et al., 2014) that provide physiological measures that track the decision process.

On the other hand, instead of looking at the evolution of the process, we can examine its final outcome. The boundary hit by the accumulation process and the hitting time are important outcomes of bounded accumulation process that relate to speed and accuracy data. The joint probability density of response time and response formalizes predictions for behavior based on the assumptions about the decision process and non-decision processes. To obtain the joint density we need to combine a model of parameter variation with a model of behavior dependent on the parameters. Let $f(t^{rt}, r \mid \boldsymbol{\theta})$ be a model of response time t^{rt} and response r , and $g(\boldsymbol{\theta} \mid \boldsymbol{\phi}, \boldsymbol{\Sigma})$ be a model of across-trial variation in parameters. Then a natural way to combine the two densities in a way that leads to predictions, and can be fit to behavioral data, is to form a mixture model (Casella and Berger, 2002).

A mixture model represents a situation when observables belong to latent subpopulations and we want to describe variation of observables in the whole population. In our case, $\boldsymbol{\theta}$ is unknown and is not of interest, so to describe variation in (t^{rt}, r) we form a mixture model

$$f(t^{rt}, r \mid \boldsymbol{\phi}, \boldsymbol{\Sigma}) = \int_{\Theta} f(t^{rt}, r \mid \boldsymbol{\theta}) g(\boldsymbol{\theta} \mid \boldsymbol{\phi}, \boldsymbol{\Sigma}) d\boldsymbol{\theta}. \quad (2.2)$$

With a mixture model we can obtain predictions for some of the features of simple choice behavior demonstrated in the Ratcliff and Rouder (1998) dataset and see how it is distinguished from the standard models with independent parameters. It would be of most interest to understand what the mixture model predicts for response time probability densities of different responses, response probabilities and relations between response times and

response probabilities.

2.3 Correlation Structure in a Benchmark Dataset

The mixture model defined in the previous section is a description of how response times and responses vary. With appropriate statistical methods, I can use the mixture model to analyze the Ratcliff and Rouder (1998) benchmark dataset. The main question is what is the correlation structure of processing parameters underlying simple choice behavior in a brightness discrimination task. I will make the simplest assumption that the correlation structure affects the decision process on each trial and it is stable across the whole experiment. The stability implies that experimental manipulations used during collection of the benchmark dataset, that is speed-accuracy instructions and brightness level, only affect means of processing parameters, but not their correlations. The simplicity of the assumption reflects ignorance about the adaptive character of the underlying decision process. Simplicity, however, also makes resultant data models maximally falsifiable.

Given that the correlation structure is stable, there are several subquestions that should guide statistical analysis to gain more information about it. The first question involves what the benchmark dataset implies about the correlation pattern? We want to know the magnitudes and directions of correlations between all the variable parameters. This question can be formalized statistically as estimation of the correlation parameters.

Another informative subquestion is how well standard models match pre-

dictions of generalized models with respect to key features of behavioral data. Graphical comparison of restricted and full models should reveal whether there is a feature of the data that discriminates between models and requires the additional assumption of correlation structure.

Lastly, we can ask whether assuming a correlation structure leads to overall improvement in predictive accuracy. This question is typically formalized with some model comparison index, which quantifies overall predictive accuracy adjusted for the complexity of each model. Comparison of the standard model to its generalizations based on a model comparison index can measure overall support for the presence of a stable correlation structure across the whole experiment.

Chapter 3

Models

With questions motivating this thesis established, I will now describe one of the standard sequential sampling models and how we might generalize it. As I discussed in 1, there are many sequential sampling models that account for performance of two-choice tasks in a similar manner. I propose to use a model developed by Ratcliff and colleagues that belongs to a class of continuous time, continuous evidence accumulators with a relative stopping rule.

The Ratcliff diffusion model is a variant of the Wiener diffusion model (Laming, 1968; Link and Heath, 1975; Ratcliff, 1978; Ratcliff and Rouder, 1998; Ratcliff and Tuerlinckx, 2002). I concentrate on this model because of its long development (Ratcliff, 1978; Ratcliff and Rouder, 1998; Ratcliff, 2002; Ratcliff and Tuerlinckx, 2002; Ratcliff, 2013), and successful testing with a wide range of experimental data (Ratcliff and McKoon, 2008; Wagenmakers, 2009). In the following section I present a mathematical description of the

Ratcliff model, its psychological interpretation, and show what features of performance data the model can and cannot predict. The model presentation follows with an introduction of the notion of copulas and how they may be used to construct flexible multivariate distributions. Last, I will propose two multivariate distributions of parameter variation that effectively generalize the Ratcliff model.

3.1 Ratcliff Diffusion Model

The Ratcliff diffusion model is arguably the most successful sequential sampling model. Underlying the Ratcliff model is a drift-diffusion Wiener process, a special case of a continuous time, continuous state space diffusion process that begins at some starting point and evolves until it reaches one of two absorbing boundaries (Ross, 2014; Karlin and Taylor, 1975, 1981). Formally, a Wiener process is an uncountable, ordered collection of univariate random variables

$$\{x(t); t \in T\} \tag{3.1}$$

with the index set $T \in [0, \infty)$ (Karlin and Taylor, 1981; Smith, 2000). Given two absorbing boundaries at 0 and $\alpha > 0$, and a starting point $\zeta \in (0, \alpha)$, the state space \mathcal{S} of $x(t)$ is the interval $[0, \alpha]$. A realization of the Wiener process, called a sample path, is a real-valued univariate function that maps the index set into the state space. In other words, a sample path is a particular sequence of states through which the Wiener process evolves until it hits a boundary and terminates.

Evolution of the drift-diffusion Wiener process is characterized by drift and diffusion coefficients. A drift coefficient,

$$\delta = \lim_{\Delta t \rightarrow 0} \frac{1}{\Delta t} \mathbb{E} [x(t + \Delta t) - x(t) \mid x(t) = x] \in \mathbb{R}, \quad (3.2)$$

quantifies mean instantaneous change in the state of the process and a diffusion coefficient,

$$\sigma^2 = \lim_{\Delta t \rightarrow 0} \frac{1}{\Delta t} \text{Var} [x(t + \Delta t) - x(t) \mid x(t) = x] > 0, \quad (3.3)$$

quantifies mean instantaneous squared change in the state of the process.

The drift coefficient determines the direction and magnitude with which an instance of the Wiener process tends to evolve towards a boundary. Positive values bias evolution towards α and negative values towards 0, while the absolute value of δ controls how rapidly the process approaches an absorbing boundary. In contrast, the diffusion coefficient determines variability in evolution. For small values the state of the process may fluctuate a bit, but for large values the process may show large fluctuations in direction and moment-to-moment transitions.

There are several equivalent ways to complete a definition of the Wiener process assuming the above facts. One definition arises from considering the evolution of the Wiener process, formalized by the linear, first-order stochastic differential equation

$$dx(t) = \delta dt + \sigma dw(t), \quad (3.4)$$

where δ and σ are defined as above, and $x(t)$ is a standard Wiener process with points distributed as normal variates with mean 0 and variance dt .

Change in the process is independent of time and current state of the process, and is a sum of an average increment and scaled Gaussian noise. Applying a stochastic integral leads to the Wiener process $x(t)$ as a unique solution to Equation 3.4.

Another approach is to describe the initial and boundary conditions, and the properties of increments, $x(t_2) - x(t_1)$, of the process. The drift-diffusion Wiener process $x(t)$ with absorbing boundaries satisfies:

1. $x(t)$ is an almost surely continuous function,
2. $x(0) = \zeta$, where $0 < \zeta < \alpha$,
3. the value of the probability density at absorbing boundaries is $f(x(t) = 0 \mid x(0) = \zeta) = f(x(t) = \alpha \mid x(0) = \zeta) = 0$,
4. $x(t) \sim \mathcal{N}(\delta t + \zeta, \sigma^2 t)$,
5. the increments are stationary; that is for $t_1, t_2 \in T$, where $0 < t_1 < t_2$, $X(t_2) - X(t_1)$ is equal in distribution to $X(t_2 - t_1)$, and
6. for $t_1, t_2, \dots, t_n \in T$, with $t_1 < t_2 < \dots < t_n$, the increments $X(t_1), X(t_2) - X(t_1), \dots, X(t_n) - X(t_{n-1})$ are independent.

From this definition, we see that the Wiener process starts at ζ and continuously evolves until it is absorbed in one of two boundaries. If we took a sample of states over a discrete time grid, then all the increments are independent normal variates that depend only on the differences between time points, but not absolute time.

In addition to its mathematical assumptions, the Ratcliff diffusion model provides linking assumptions that map the Wiener process and its parameters into a psychological process and variables relevant to describing a decision mechanism. In line with sequential sampling approach to decision making, the drift-diffusion Wiener process models a perfect but noisy integrator of relative information. This integrator starts at some initial amount of decision-relevant information and continuously accumulates continuous information samples from a sensory representation without loss of information (Smith, 2000; Bogacz et al., 2006).

The index set $T = [0, \infty)$ describes the time passed from the beginning of the accumulation, and the state space $S = [0, \alpha]$ describes the possible amounts of information that an integrator can hold under certain conditions. The drift coefficient δ measures the average rate of information uptake by a perfect integrator, and the diffusion coefficient σ^2 describes the level of noise in the information uptake.

The boundaries of the state space, 0 and α , represent decision thresholds that quantify the amount of relative information needed to form a decision. The lower and upper thresholds could represent each of two choices, say “No” and “Yes” in the auditory signal detection experiment. The separation between boundaries, α , quantifies decision caution, with larger separation implying that more information is required before making a decision. The two thresholds force a relative stopping rule where information is relative because movement towards one threshold is movement away from the other.

The starting point ζ of the process represents the initial amount of in-

formation from which the accumulation process begins. The level of initial information represents the bias an observer may have towards one of two choices. The farther away from $\alpha/2$ the starting point ζ is, the more initial information there is for a particular choice. The initial information asymmetry suggests an alternative parameterization of the Wiener process corresponding to a biased integrator that starts accumulation with bias $\beta = \zeta/\alpha \in [0, 1]$, a proportion of the distance from 0 to α .

The drift-diffusion Wiener process model of a perfect integrator provides a natural way to predict the joint probability density of decisions and decision times in a two-choice task. A sample path of the drift-diffusion Wiener process will start at ζ (or β) and take some random time to absorb at one of the boundaries. Let $t^d > 0$ be the time to absorption and $r \in \{0, 1\}$ be the boundary passed at t^d , with 1 corresponding to boundary at α . Translating into psychological terms, the sample path is the evolution of the information integrator to a decision r during the duration of decision time t^d .

The primary prediction of the Ratcliff diffusion model is the joint variation of decisions and decision times. The component of the joint probability density function when $r = 1$

$$f(t^d, 1 \mid \alpha, \delta, \zeta, \sigma^2) = \frac{\pi\sigma^2}{\alpha^2} \exp\left(\frac{\delta\zeta}{\sigma^2} - \frac{\delta^2 t^d}{2\sigma^2}\right) \sum_{k=1}^{\infty} k \exp\left(\frac{-k^2 \pi^2 \sigma^2 t^d}{2\alpha^2}\right) \sin\left(\frac{k\pi(\alpha - \zeta)}{\alpha}\right) \quad (3.5)$$

and the component for the opposite response is obtained by replacing δ with $-\delta$ and $\alpha - \zeta$ with ζ . The form of Equation 3.5 suggests that distributional prediction of the Wiener process model suffers from lack of identifiability A

model is identifiable when for any nonequivalent parameters $\theta_1, \theta_2 \in \Theta$, the resulting probability density functions are unique, that is $f(x | \theta_1) \neq f(x | \theta_2)$ (Casella and Berger, 2002). As a result, identifiability of a model creates the possibility of precise estimation of true parameters with an infinite sample size. For Equation 3.5, assuming that $\{\alpha, \delta, \zeta, \sigma^2\}$ are free to vary, if one multiplies all parameters of the model by a factor c , then the predictions do not change. Thus, parameter recovery is unique upto a scaling operation.

For the Wiener model, resolving this lack of identifiability requires fixing one of the parameters to an arbitrary value. Typically, application of the Wiener model relied on fixing the diffusion coefficient σ^2 to 0.1 or 1 across experimental conditions (Ratcliff, 1978; Ratcliff and Tuerlinckx, 2002; Ratcliff and McKoon, 2008). The stability of noise is a theoretical assumption, and ultimately should be tested against a particular experimental situation. Given that this thesis concentrates on the Ratcliff and Rouder (1998) data, which was generated without manipulations aimed at changing noisiness of the integrator, I will assume that $\sigma^2 = 1$, across all experimental conditions, to resolve this identifiability issue.

The model specified so far allows predictions only about t^d , which represents the unobservable decision time. To fully connect the Wiener process to behavioral data, the model requires assumptions about the perceptual and motor processes surrounding the decision stage. A sufficient adjustment is to relate the decision time to the response time because the response is assumed to reflect the decision without error, or that motor processing exactly reflects which evidence bound was reached.

An additive decomposition of two-choice tasks implies that before information sampling can begin, the sensory stimulus has to be transduced and its features encoded. When the accumulation reaches one of the thresholds, the observer has to organize and execute a motor response for his or her decision to become observable. The mechanistic details of such perceptual and motor processing are typically not specified when specifying a sequential sampling model (but see Smith and Ratcliff (2009); Brown et al. (2008)), but their time can be captured by a single parameter τ^{nd} . Then, response time

$$t^{rt} = t^d + \tau^{nd} \quad (3.6)$$

where t^d is the decision time and τ^{nd} is the combined perceptual and motor processing (non-decision) time. If τ^{nd} is a constant, the joint probability density of response times and responses is a shifted density of decision times and decisions

$$f(t^{rt} - \tau^{nd}, r \mid \alpha, \delta, \zeta, \sigma^2) = f(t^{rt}, r \mid \alpha, \delta, \zeta, \sigma^2, \tau^{nd}), \quad (3.7)$$

with support over $t^{rt} \in [\tau^{nd}, \infty)$ and $r \in \{0, 1\}$.

Comparison of the Wiener model's predictions with experimental data shows that, for some experimental paradigms, the model can accurately predict several qualitative and quantitative features of behavioral data. First, large variation in response times, as shown in Figure 1.1, arises from within-trial variability in accumulation. Variability follows because sometimes the integrator uptakes information rapidly resulting in a fast response time. On other trials, the integrator may oscillate for a while around the initial state before terminating, leading to a slow response time.

Second, noisy within-trial accumulation accounts for errors. To model accuracy one has to assume that the starting point $\zeta = \alpha/2$, the drift rate $\delta \geq 0$ and boundaries represent correct (upper) and error (lower) responses. The starting point restriction reflects an assumption that there is no bias towards correct or error responses, and the non-negative drift rate represents the assumption that stimulus information restricts performance to be between chance and full accuracy. With these assumptions, noisy accumulation sometimes results in the integrator terminating at the lower bound, which causes an error.

Third, the decision model based on the Wiener process can predict the positive skew of response time densities observed in Figure 1.1. The Wiener process tends to move directly towards one of the bounds, and only sometimes oscillates between the bounds or reverses its course, which predicts that response time probability densities for reaching $r = 1$ and $r = 0$ will have a positively skewed shape.

Fourth, the model can explain the speed-accuracy tradeoff. Experimentally, the speed-accuracy tradeoff is obtained by manipulating performance instructions. The Wiener model can predict the trade off, like the one in Figure 1.2, by adjusting the threshold parameter, α , or, in psychological terms, using instructions to manipulate a participant's level of caution. Assuming an unbiased starting point, $\zeta = \alpha/2$, and $r = 1$ representing the correct response, then increasing α will result in the integrator taking a longer time to produce the correct response. The farther distance from ζ to 0 will also increase accuracy by allowing some of the accumulationg paths initially mov-

ing towards the error bound to reverse their trajectory towards the correct bound. However, recall that there are multiple studies either of physiological or cognitive modeling type that conclude that other parameters need to simultaneously adjust with α to full explain effects of instructions (Rinkenauer et al., 2004; Vandekerckhove et al., 2011; Rae et al., 2014).

Fifth, the Wiener model predicts that with higher accuracy participant’s correct response times will tend to be faster. This pattern is observed for both speed and accuracy instructions, as shown in Figure 1.2. When instructions are fixed, only the stimulus quality is manipulated. Intuitively, with better stimulus information the task becomes easier, so the accuracy goes up and processing time is faster. The observed relation is predicted by assuming that the quality of stimulus, such as the ratio of white to black pixels, modulates the drift rate. When $r = 1$ is the correct response, the Wiener model captures the joint effect because increasing the drift rate forces accumulation to move more rapidly and more frequently towards the upper bound.

The five features captured by the Wiener model are only a fraction of the full collection of regularities found in the benchmark and other datasets. The model is theoretically inadequate to account for asymmetric relations between correct and error response times, additional variation in the left tail of the response time density and auto-correlation of response times and responses.

Motivated by some failures of the Wiener model, the Ratcliff model adds additional structure to the Wiener process that helps to account for asymmetry between correct and error response times as well as additional left

tail variation (Ratcliff, 1978; Ratcliff and Rouder, 1998). The solution to these empirical failures is to propose additional noise in cognitive processing, which manifests itself in random fluctuations of parameters across trials. Specifically, the Ratcliff model assumes that there is trial-to-trial variation in stimulus quality, initial state of the integrator and non-decision time.

Mathematically, the idea of processing variation can be modeled with the drift-diffusion Wiener process with parameters varying across trials according to some multivariate distribution. Thus, a sample of response times arises from a mixture of Wiener processes evolving under different combinations of parameters. The full Ratcliff diffusion model assumes that the starting point ζ , non-decision time τ^{nd} and drift rate δ are random variables with probability densities

$$\begin{aligned}\zeta &\sim \mathcal{U}(\lambda - \frac{\gamma}{2}, \lambda + \frac{\gamma}{2}), \\ \tau^{nd} &\sim \mathcal{U}(\chi - \frac{\phi}{2}, \chi + \frac{\phi}{2}), \text{ and} \\ \delta &\sim \mathcal{N}(\nu, \eta^2),\end{aligned}\tag{3.8}$$

where $\lambda > 0$ and $\chi > 0$ are mean parameters, $\gamma > 0$ and $\phi > 0$ are range parameters, $\nu \in \mathbb{R}$ is a mean parameter and $\eta > 0$ is a variance parameter. To make sure that the starting point is within the state space and non-decision time is positive, the parameters satisfy $0 < \lambda - \frac{\gamma}{2} < \lambda + \frac{\gamma}{2} < \alpha$ and $0 < \chi - \frac{\phi}{2}$.

The distributional assumptions for ζ , τ^{nd} and δ were made for theoretical reasons (Ratcliff, 1978; Ratcliff and Rouder, 1998; Ratcliff and Tuerlinckx, 2002). For example, normally distributed drift represents the assumption of noisy stimulus effects such that, for nominally identical stimuli, the internal

representation varies due to a multitude of independent random influences, which gives rise to overall normal variation by the central limit theorem (Vickers, 1979; Casella and Berger, 2002). The starting point variability may reflect momentary biases and residual evidence across trials, and variability in non-decision time reflects noisiness of the peripheral nervous system. In addition to these theoretical reasons, the distributional assumptions simplify parameter estimation, which traditionally is formulated as optimizing a measure of goodness-of-fit (Tuerlinckx, 2004; Vandekerckhove and Tuerlinckx, 2007; Voss and Voss, 2007).

Given that many other assumptions can be made, a natural concern is the sensitivity of model predictions and substantive conclusions to assumptions of across-trial variability. Ratcliff (2013) demonstrated that the model fit is robust with respect to a wider class of distributions. For a wide range of parameters, swapping in a beta for the normal distribution of drift, a beta for the uniform starting point, a normal or exponential for the uniform nondecision time, one at a time, resulted in accurate recovery of parameters and the same substantive conclusions. Except when parameter values were extreme, the assumption of exponentially-distributed nondecision time, or normally-distributed nondecision time with small boundary α , the study showed that empirical success of the model is moderately insensitive to across-trial variability assumptions.

The Ratcliff assumptions imply that the first passage density is a statis-

tical mixture model with joint density function

$$f(t^{rt}, r \mid \alpha, \lambda, \gamma, \nu, \eta, \chi, \phi) = \int \int \int_V f(t^{rt}, r \mid \alpha, \zeta, \delta, \tau^{nd}) f(\zeta \mid \lambda, \gamma) f(\delta \mid \nu, \eta) f(\tau^{nd} \mid \chi, \phi) d\zeta d\delta d\tau^{nd}, \quad (3.9)$$

where $r = 0, 1$ and the individual densities have been specified above.

This probability density determines all the falsifiable predictions of the Ratcliff diffusion model for two-choice behavioral data and can accurately describe the five features of performance data mentioned above, as well as account for two additional features (Ratcliff and Tuerlinckx, 2002; Ratcliff, 2002).

The first additional data feature the Ratcliff model can handle is occasional asymmetry between correct and error mean response times (Ratcliff and McKoon, 2008; Wagenmakers, 2009). An asymmetry reflects correlation between response times and responses. There are three basic patterns that can arise: slower errors, faster errors, and a cross-over pattern where for easy stimuli errors are faster and difficult stimuli errors are slower. The benchmark data shows a cross-over pattern, as shown in Figure 1.2. These patterns are very diagnostic, and have long been a target for model development (Laming, 1968; Ratcliff, 1978; Vickers, 1979; Ratcliff and Rouder, 1998).

The Ratcliff model can handle all three asymmetry patterns with the across-trial variation assumptions. The slow errors arise only from assuming that the drift rate δ follows a normal distribution. How drift variation produces slow errors can be understood by considering that the mean cor-

rect and error response times for the mixture model are weighted averages of correct and error response times with respect to variation in drift rate values.

Assume that t_c^{rt} is the correct response time and t_e^{rt} is the error response time. Then, assuming only the drift rate varies, the correct mean of the mixture model

$$\mathbb{E} [T_c^{rt}] = \mathbb{E}_\delta [\mathbb{E} [T_c^{rt} | \delta]] = \int_{\mathbb{R}} \mathbb{E} [T_c^{rt} | \delta] f(\delta | \nu, \eta^2) d\delta, \quad (3.10)$$

and similarly the error mean of the mixture model

$$\mathbb{E} [T_e^{rt}] = \mathbb{E}_\delta [\mathbb{E} [T_e^{rt} | \delta]] = \int_{\mathbb{R}} \mathbb{E} [T_e^{rt} | \delta] f(\delta | \nu, \eta^2) d\delta, \quad (3.11)$$

where $f(\delta)$ is a normal density with mean ν and variance η . Suppose stimulus favours $r = 1$ and ν is positive, so the mass of the $f(\delta)$ is heavily distributed over positive numbers. For a positive drift rate, the error mean is slower than the correct mean. Thus, averaging mean correct and error response times with respect to $f(\delta)$ will result in upweighing fast correct means and slow error means. Hence, the overall mixture pattern would be slower errors.

Fast errors can be incorporated into a Ratcliff diffusion model by assuming variability in the starting point ζ . With drift rate fixed to a positive value, most Wiener processes will tend to terminate at the correct bound, so for errors to occur the process has to rapidly evolve towards the lower bound. Thus, mean error response times for starting values closer to the lower bound will tend to be faster than mean correct response times for symmetric starting values. When equilly weighted by the uniform density, the mixed means will result in faster errors.

Combining variability in the drift rate and starting point in the Ratcliff model can produce all three patterns. The switch between different patterns occurs from relative differences in the standard deviations of the drift rate and the starting point, or from the relative means. When variability of the drift rate is sufficiently higher than variability of the starting point, the pattern of slow errors arises. Reversing the relation will result in fast errors. For relatively close levels of variability, by shifting the mean of the drift rate from high to low, performance moves from fast errors to slow errors, as evidenced in Figure 1.2.

The second additional pattern the Ratcliff model predicts is extra variability in the lower tail not predicted by the Wiener model. The pattern can be quantified by examining the 0.1 quantile. Assuming trial-to-trial variability in the non-decision time τ^{nd} allows the Ratcliff model to predict this empirical pattern. The reason is that the non-decision parameter sets the lower bound on the fastest response time, so the effect of across-trial variation is strongest in the left tail. Psychologically, fast response times correspond to the decision process being around its peak speed, so most variability arises from perceptual and motor processing.

Even with all the variability assumptions, the Ratcliff model still cannot capture two important features of behavioral data. First, it is typical to observe extremely fast or slow response times within a series of trials, and second, serial dependencies among responses and response times are ubiquitous, as evidenced for the benchmark data in the Figure 1.3. Several studies have proposed ways to model these features (Ratcliff and Tuerlinckx, 2002;

Vandekerckhove and Tuerlinckx, 2007; Gao et al., 2009; Craigmile et al., 2010; Jones et al., 2013). The auto-correlation in the response times can be captured by assuming, that next trial’s parameters somehow depend on the current or even past trial’s parameters. Learning and information carry-over are two plausible psychological process that could underlie such auto-correlation. In contrast, the extreme response times can be accounted by fast guesses or delayed decisions, and modeled with a mixture of Wiener processes, delayed Wiener processes and a guessing process.

Overall, the Ratcliff diffusion model accounts accurately for many features of behavioral data coming from two-choice tasks. For this reason I will use it as a basis for investigating parameter correlations. One of the questions motivating this thesis is how to model correlations between parameters. As discussed above, I will take a descriptive approach that requires a multivariate density function to describe parameter variation. The next section introduces a method for defining new multivariate densities based on the concept of copulas.

3.2 Copula-based Multivariate Distributions

Generalizing existing simple decision models to have correlated parameters requires a multivariate distribution that captures variances and covariances of the parameters. The problem of finding a candidate distribution is both underconstrained, given the lack of information about the covariance structure, and restricted, because the parameters do not have the same scale. Hence,

the method of distribution construction should allow combining an arbitrary covariance structure and arbitrary univariate marginal distributions.

For simplicity, consider the example of standard signal detection theory. Recall that the model of hits and false alarms is controlled by three free parameters: $\mu_s \in \mathbb{R}$, $\sigma_s > 0$, and $\tau \in \mathbb{R}$. The usual assumptions say that these parameters do not vary across trials and have no mutual restrictions. Suppose in a paradigm where stimulus level changes across trials, such as in the Ratcliff and Rouder (1998) brightness discrimination experiment, representational parameters μ_s and σ_s could vary. This variation could be such that, when μ_s increases in response to a stronger stimulus, σ_s tends to decrease. This negative correlation corresponds to a representation process with a noise suppression property. If we are interested in testing such a representational hypothesis, then we need to formalize parameter variation in a bivariate distribution for μ_s and σ_s . The obstacles to such a formalization, similar to simple decision models, are that the exact nature of the relation between μ_s and σ_s is not known and they have different scales.

The framework of copulas, developed in probability theory over the last 60 years and recently imported for data analysis in a variety of fields, is one way to handle obstacles to modeling across-trial variability (Sklar, 1959; Joe, 1997; Nelsen, 2007; Berkes et al., 2008). A copula C is a probability distribution function with a unit hypercube support, so for some parameter $\boldsymbol{\theta} \in \Theta \subseteq \mathbb{R}^k$ and random variables $u_i \in [0, 1]$, $C(u_1, u_2, \dots, u_p \mid \boldsymbol{\theta}) : [0, 1]^p \mapsto [0, 1]$. The dimensions of a copula may be independent or dependent, but each dimension is a continuous uniform random variable $u_i \sim \mathcal{U}(0, 1)$, $i =$

$1, 2, \dots, p$, which has no free parameters. Thus, parameter $\boldsymbol{\theta}$ only provides information about joint variation, but nothing about marginal variation.

To motivate the relevance of copulas to describing across-trial variability, let's consider a problem of sampling hits and false alarms from a signal detection model with variable μ_s and σ_s for a sequence of variable signal stimuli. Formally, for a sequence of n signal stimuli, responses r are generated hierarchically according to

$$\begin{aligned} (\mu_s^{(i)}, \sigma_s^{(i)})^T &\sim F(\boldsymbol{\theta}) \\ r^{(i)} &\sim f(r \mid \mu_s^{(i)}, \sigma_s^{(i)}, \tau), \end{aligned}$$

where F is an unknown distribution and f is a Bernoulli probability mass function. One defining feature of F is how μ_s and σ_s covary. The properties of copulas stated above suggest that copulas are a complete description of statistical dependencies between standardized variables. A hypothesis of negative association between μ_s and σ_s , whether linear or nonlinear form, could then be expressed as a copula. However, a sample from a copula is not enough to get a sample of (μ_s, σ_s) because their scales differ from unit interval.

Given a sample from a copula, we need to somehow map the standardized $[0, 1]$ dimensions of a copula onto \mathbb{R} and $(0, \infty)$ scales, and change continuous uniform distributions into desired marginal distributions, say normal and gamma. A useful theorem from probability theory, probability integral transform theorem, says that a univariate random variable $x \sim F$ relates to a uniform random variable via an equation $x = F^{-1}(u)$ (Casella and Berger,

2002). As a consequence of this fact, if we had a random sample of standard uniform variables $u_1, u_2, \dots, u_n \sim \mathcal{U}(0, 1)$ and arbitrary distribution F , then we could get a random sample of x by applying $x = F^{-1}(u)$.

Going back to the signal detection example, assume we want μ_s to be normally distributed and σ_s to be gamma distributed. Then, using the probability integral transform theorem, we could transform a random sample from a copula C , encoding a desired form of negative correlation between standardized variables, into μ_s and σ_s^2 by applying normal and gamma marginal distributions to the standardized variables. A random sample of μ_s and σ_s could then be used to obtain predictions of hits and false alarms for n signal stimuli according to the hierarchical scheme.

The prediction of responses under variable signal stimuli suggests that with a copula encoding negative associations, and normal and gamma marginal distributions of parameters, we could obtain a joint distribution of parameters that encodes our hypothesis about across-trial variability. The basis for the connection between copulas and multivariate distributions derives from Sklar's 1959 two-part representation theorem:

Theorem. *Let F be a multivariate distribution with continuous marginal distributions F_1, F_2, \dots, F_p , then there exists a unique copula C such that*

$$F(x_1, x_2, \dots, x_p) = C(F_1(x_1), F_2(x_2), \dots, F_p(x_p)). \quad (3.12)$$

Conversely, given a copula C and marginal univariate distributions F_1, F_2, \dots, F_p , F is a multivariate distribution.

A corollary of Sklar's theorem, whose importance will be explained below, is

Corollary 1. *Let F be a p -dimensional distribution function with continuous quantile functions $F_1^{-1}, F_2^{-1}, \dots, F_p^{-1}$, then for every $(u_1, u_2, \dots, u_p) \in [0, 1]^p$ a copula*

$$C(u_1, u_2, \dots, u_p) = F(F_1^{-1}(u_1), F_2^{-1}(u_2), \dots, F_p^{-1}(u_p)). \quad (3.13)$$

The forward implication of Sklar's theorem means that every multivariate distribution can be rewritten in terms of a copula and univariate marginals. One way to think about multivariate distributions is that they simultaneously describe joint and marginal variation of a random vector. The right hand side of Sklar's equation is a combination of a copula and marginal distributions. On its own, a copula carries no information about the range of each dimension of a random vector or the scale, shape and location of each variable. Hence, a copula only describes joint variation. On the other hand, a univariate distribution function defines the range of a variable and other characteristics through its parameters, thus specifying marginal variation.

Let's examine how the forward implication applies to the joint distribution of hits and false alarms of the standard signal detection model by deriving its copula. By definition, the joint distribution function of the number of hits x_h and false alarms x_{fa}

$$F(x_h, x_{fa}) = P\{x_h \leq k_h, x_{fa} \leq k_{fa}\}, \quad (3.14)$$

where $k_h = 0, 1, \dots, N_h$, $k_{fa} = 0, 1, \dots, N_{fa}$. We can further rewrite the joint

distribution as

$$\sum_{x_h=0}^{k_h} \sum_{x_{fa}=0}^{k_{fa}} p(x_h, x_{fa}) = \sum_{x_h=0}^{k_h} \sum_{x_{fa}=0}^{k_{fa}} p_h(x_h) p_{fa}(x_{fa}) = \sum_{x_h=0}^{k_h} p_h(x_h) \sum_{x_{fa}=0}^{k_{fa}} p_{fa}(x_{fa}), \quad (3.15)$$

which simplifies to

$$F_h(x_h) F_{fa}(x_{fa}). \quad (3.16)$$

By Sklar's theorem, the copula for the signal detection model is

$$C(u_1, u_2) = u_1 u_2, \quad (3.17)$$

where $u_1, u_2 \in [0, 1]$. The functional form of C says that joint variation between hits and false alarms can be described as a uniform distribution over a unit square, implying that knowing the number of hits tells us nothing about the number of false alarms. This follows from the standard assumption of independent responses in signal detection theory and is reflected in the special case of a n-dimensional independence copula.

The derivation of the independence copula from a bivariate distribution of hits and false alarms shows how Sklar's theorem is also a statement about how copulas can be obtained from known multivariate distributions. Once a copula is derived, we can study joint variation among random variables while ignoring that, marginally, they may be binomial or otherwise. In contrast, the converse implication of Sklar's theorem takes us from a copula and marginal distributions to a related multivariate distribution. The practical consequence is that if we can specify joint variation with a copula and marginal variation with some univariate distributions, then Sklar's theorem gives us a way to define new multivariate distributions.

To demonstrate the usefulness of the converse implication, consider two alternatives of the standard signal detection model that incorporate variability in parameters of a signal representation. To specify two models we need copulas and marginal distributions to characterize variation in μ_s and σ_s . As a first candidate to describe negative correlation between the parameters, consider a Frank copula

$$C_F(u_1, u_2) = -\frac{1}{\theta_F} \ln \left(1 + \frac{(\exp(-\theta_F u_1) - 1)(\exp(-\theta_F u_2) - 1)}{(\exp(-\theta_F) - 1)} \right) \quad (3.18)$$

where $\theta_F \in \mathbb{R} \setminus \{0\}$ is a correlation parameter. A sample from a Frank copula can be seen in Figure 3.1. The characterizing feature of a Frank copula is that u_1 and u_2 covary over a rectangular-like contour.

An alternative to a Frank copula is a t copula

$$C_t(u_1, u_2) = \int_{-\infty}^{t_\nu^{-1}(u_1)} \int_{-\infty}^{t_\nu^{-1}(u_2)} \frac{1}{2\pi(1-\theta_t^2)^{1/2}} \left(1 + \frac{x^2 - 2\theta_t xy + y^2}{\nu(1-\theta_t)} \right)^{-(\nu+2)/2} dx dy, \quad (3.19)$$

where $\nu > 0$ is a degrees of freedom, $\theta_t \in [-1, 1]$ is a correlation parameter, $t_\nu^{-1}(\cdot)$ is a t quantile function, $x \in \mathbb{R}$ and $y \in \mathbb{R}$. Note that the t copula is obtained by applying Corollary 1 to the bivariate t distribution function, revealing the importance of the corollary for deriving new copulas. A sample from a t copula is shown in Figure 3.1. The t sample shows that dependence between u_1 and u_2 has an elliptical contour with extended tails.

By picking Frank and t copulas we can now describe joint variation in μ_s and σ_s . To complete defining models of across-trial variation we need two univariate distributions. For μ_s I will use a non-central t distribution F_t with degrees of freedom $\omega > 0$ and a shape parameter $\lambda \in \mathbb{R}$ that introduces a




Figure 3.1: The top two panels show samples from a Frank copula and a t copula. Bivariate distributions resulting from applying a t and a gamma transformation to copula samples are in the bottom panel. Parameter settings: $\theta_F = -15, \theta_t = -0.95, \omega = 2, \nu = 3, \lambda = 1.5, \alpha = 3, \beta = 6$. Parameters result in approximately the same Kendall's measure of correlation.

skew. The support of t covers the real number line, so it satisfies the range requirement for μ_s . As for σ_s , I will assign a gamma distribution F_g , which has support over the positive real numbers, with a shape parameter $\alpha > 0$ and a scale parameter $\beta > 0$.

The two copulas can now be combined with the marginal distributions according to Sklar's theorem to give us two distributions of μ_s and σ_s . Model M_1 takes the form of

$$F_1(\mu_s, \sigma_s \mid \theta_F, \omega, \lambda, \alpha, \beta) = C_F(F_t^{-1}(u_1 \mid \omega, \lambda), F_g^{-1}(u_2 \mid \alpha, \beta) \mid \theta_F), \quad (3.20)$$

and model M_2 is

$$F_2(\mu_s, \sigma_s \mid \theta_t, \nu, \omega, \lambda, \alpha, \beta) = C_t(F_t^{-1}(u_1 \mid \omega, \lambda), F_g^{-1}(u_2 \mid \alpha, \beta) \mid \theta_t, \nu). \quad (3.21)$$

Samples from M_1 and M_2 are shown in Figure 3.1. M_1 , on the bottom left, induces a negative, monotonically decreasing relation that shows smooth co-variation between μ_s and σ_s . M_2 , on the bottom right, also induces a negative correlation, but has an abrupt shift in the “slope” after $\mu_s \approx 2$, suggesting two different states that a representation may occupy. Overall, Frank and t copulas result in two bivariate distributions of the signal representation properties that have a noise suppressing feature.

Using the two models we could examine effects of correlated parameters on hits and false alarms. Before showing how to integrate the copula-based distributions with the bivariate binomial mass function, I will present another useful corollary of Sklar's theorem.

Corollary 2. *For a differentiable p -dimensional distribution function with continuous quantile functions $F_1^{-1}, F_2^{-1}, \dots, F_p^{-1}$, the probability density function*

$$\begin{aligned}
f(x_1, x_2, \dots, x_p) &= \frac{\partial^p F(x_1, x_2, \dots, x_p)}{\partial x_1 \partial x_2 \cdots \partial x_p} \\
&= \frac{\partial^p C(F_1(x_1), F_2(x_2), \dots, F_p(x_p))}{\partial x_1 \partial x_2 \cdots \partial x_p} \\
&= c(F_1(x_1), F_2(x_2), \dots, F_p(x_p)) \prod_{i=1}^p f(x_i), \quad (3.22)
\end{aligned}$$

where $c(F_1(x_1), F_2(x_2), \dots, F_p(x_p))$ is a probability density associated with a copula, and $f(x_i)$ are probability densities of each dimension.

One important consequence of Corollary 2 is that we can derive a copula density $c(F_1(x_1), F_2(x_2), \dots, F_p(x_p))$ by dividing the joint density $f(x_1, x_2, \dots, x_p)$ by the product of marginal densities $\prod_{i=1}^p f(x_i)$, revealing again that a copula is the factor that adds dependence between random variables. From the practical side, the implication is that we can define a new multivariate distribution by going directly to the probability density functions. This becomes useful when marginal distribution functions have no closed form (e.g. t distribution).

Going back to the signal detection example, we can now obtain a functional form for hits and false alarms under parameter variability. Let's concentrate on M_2 as a model of μ_s and σ_s . By Corollary 2, the density of M_2 is a product of a t copula density c with a noncentral t density f_t and a gamma

density f_g , so

$$f(\mu_s, \sigma_s \mid \theta_t, \omega, \lambda, \nu, \alpha, \beta) = c(F_t(\mu_s \mid \lambda, \nu), F_g(\sigma_s \mid \alpha, \beta) \mid \lambda) \times f_t(\mu_s \mid \lambda, \nu) f_g(\sigma_s \mid \alpha, \beta). \quad (3.23)$$

Then, in an experiment where we observe N_s signal trials and N_{fa} trials, with assumption of variable μ_s and σ_s according to M_2 , the probability mass function for hits and false alarms is

$$\int_{-\infty}^{\infty} \int_0^{\infty} p(y_h \mid \mu_s, \sigma_s, \tau, N_s) p(y_{fa} \mid \tau, N_{fa}) f(\mu_s, \sigma_s \mid \theta_t, \omega, \lambda, \nu, \alpha, \beta) d\mu_s d\sigma_s = p(y_h \mid \theta_t, \omega, \lambda, \nu, \alpha, \beta, \tau, N_s) p(y_{fa} \mid \tau, N_{fa}) \quad (3.24)$$

With two alternative models of hits and false alarms at hand, we could now fit real data and test for the assumption of a dynamic representation of signals. It would be of particular interest to obtain copula parameters. A visual summary of the best model could be captured by a graph of a copula to reveal what sort of joint variation is consistent with data.

Estimated copula parameters also provide a numerical summary of inference, but they may be more difficult to interpret and it would be beneficial to map them onto commonly used scalar measures of dependence. From Sklar's result, it turns out that copulas provide a natural way for deriving several important measures of dependence between random variables, regardless of univariate marginals.

One way to connect copula parameters to common measures of dependence may be through Pearson's scalar measure of linear dependence

$$\psi_P = \frac{\text{Cov}[x, y]}{\sqrt{\text{Var}[x] \text{Var}[y]}}, \quad (3.25)$$

but it turns out not to work. It is appropriate when samples come from a normal or a t distribution (Casella and Berger, 2002). However, Pearson's coefficient cannot be rewritten strictly in terms of a copula, and when the joint distribution is t, the coefficient may be undefined because moments may not be defined at lower degrees of freedom. Another downside is that marginal distributions of x and y have influence on the Pearson's measure through the Var operator, revealing that it is partially influenced by marginal variation.

As an alternative, Kendall's coefficient, which measures the amount of monotonic dependence between two random variables, can be defined strictly in terms of a bivariate copula (Joe, 1997; Nelsen, 2007). For a bivariate random vector $(x, y)^T$ drawn from a copula $C(\cdot)$ (with probability density $c(\cdot)$), the Kendall's coefficient equals

$$\begin{aligned} & 4 \iint_{[0,1]^2} C(u, v) c(u, v) du dv - 1 \\ &= 4\mathbb{E}[C(u, v)] - 1. \end{aligned} \tag{3.26}$$

The advantages of Kendall's coefficient is that it extends to monotonic relations between variables, does not depend on marginal distributions and provides a familiar scale for comparing different copula-based models. For example, in the generalized signal detection case, with a model of hits and false alarms based on a t copula, Kendall's coefficient is

$$\rho_k = \frac{2}{\pi} \arcsin(\theta_t), \tag{3.27}$$

which is readily computable.

In addition to general dependence, the tail dependence is a second scalar measure that can help to quantitatively characterize relations between random variables in a simple form. This quantity considers dependence imposed by a copula only in the lower-left or upper-right quadrants. The tail dependence is defined as an asymptotic probability that x_2 exceeds some quantile q given that x_1 exceeds q , corresponding to co-occurrence of variables in the extremes. The lower-left tail dependence is

$$\lambda_L = \lim_{q \rightarrow 0} P(x_2 \leq F_2^{-1}(q) \mid x_1 \leq F_1^{-1}(q)), \quad (3.28)$$

and the upper-right tail dependence is

$$\lambda_R = \lim_{q \rightarrow 1} P(x_2 > F_2^{-1}(q) \mid x_1 > F_1^{-1}(q)). \quad (3.29)$$

Similar to the Kendall's coefficient, the tail dependence can be defined strictly in terms of a copula $C(\cdot)$. For the lower-left tail, the dependence

$$\lambda_L = \lim_{q \rightarrow 0^+} \frac{C(q, q)}{q}, \quad (3.30)$$

and for the upper-right tail

$$\lambda_U = \lim_{q \rightarrow 1^-} \frac{1 - 2q + C(q, q)}{q}. \quad (3.31)$$

For example, for t and Frank copulas used in generalizing the standard signal detection model, the former has tail dependence and the latter has no tail dependence. Tail dependence can be seen in Figure 3.1, especially after applying the quantile transformations, which result in rounded tails for the Frank-based model and sharp tails in the t-based model. In this context,

tail dependence could correspond to a representational processes where noise remains tightly bound to a minimal value at high magnitudes of signal.

In conclusion, copulas provide a useful and powerful way of defining new multivariate distributions with various patterns of dependencies which can be easily summarized with scalar quantities. The signal detection model generalization examined in this section serves as a direct example of how models of simple choice behavior can be generalized to have dependent across-trial variation in their parameters. In the next section, I apply the copula-based approach to propose two multivariate distributions for parameter variation that can be used to generalize the Ratcliff diffusion model.

3.3 Proposed Generalizations

Recall that one of the questions motivating this thesis is developing a model of across-trial variability of parameters of a sequential sampling model. Following a descriptive approach I will present two new multivariate distributions that we can use to generalize the Ratcliff diffusion model. Resulting mixture models will set the stage for answering the other questions guiding this thesis.

I assume the Wiener process as a description of the decision process underlying simple choice behavior. To avoid the inconvenient domain of the starting point, I will work with the bias parameterization of the Wiener process where $\beta = \zeta/\alpha$, instead of ζ , fixes the starting value of accumulation. Three parameters are assumed to vary across trials: drift rate, bias and

non-decision time. Our problem is to find some multivariate densities that describe the behavior of the three parameters, each of which falls on a different scale and has a different univariate distribution. Sklar’s theorem tells us that to describe joint variation of the three parameters we need a copula and marginal probability density functions for each dimension.

I start with copulas. When searching for copulas I concentrated on families that were flexible and computationally tractable. I interpret flexibility as the ability of a copula to produce different correlations for each pair of parameters ranging from full negative to full positive dependence. Basically, we need copulas that can encode as many different types of correlation structures as possible. Within the set of existing copulas, many copulas, regardless of the dimension, are governed by a single parameter that induces a particular kind of dependence (Joe, 1997; Nelsen, 2007). For example, Archimedean copulas are controlled by a single parameter and encode the same lower tail, upper tail, or all-positive dependence for all pairs of variables. The Frank copula presented in the previous section is an example. Archimedean and other restricted copulas are inadequate because there is no prior information available about the correlation structure of parameters, so it is sensible to start modeling with most flexible copulas.

The elliptical class of copulas satisfies the flexibility and tractability criterion. Elliptical copulas arise from n -dimensional elliptical distributions that can be defined in terms of their probability density

$$f(\mathbf{x} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma}) = |\boldsymbol{\Sigma}|^{-1} \mathbf{g}\left((\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})\right), \quad (3.32)$$

where $|\cdot|$ is a determinant, g is a nonnegative function on the positive real numbers, $\boldsymbol{\mu} \in \mathbb{R}$ is a location parameter, and $\boldsymbol{\Sigma}$ is a positive definite matrix of scale and correlation parameters (Joe, 1997; Nelsen, 2007). The term “elliptical” refers to the quadratic form (the argument of $g(\cdot)$) that causes iso-density contours to take the shape of a hyperellipsoid. Specifically, I propose using normal and t copulas from the elliptical class.

Consider a normal copula first. Taking the three dimensional standard normal distribution function $\Phi_3(\mathbf{x} \mid \mathbf{P}_n)$, $\mathbf{x} \in \mathbb{R}^3$, with a positive definite correlation matrix

$$\mathbf{P}_n = \begin{bmatrix} 1 & \rho_{1,2} & \rho_{1,3} \\ \rho_{1,2} & 1 & \rho_{2,3} \\ \rho_{1,3} & \rho_{2,3} & 1 \end{bmatrix},$$

and the standard normal quantile function $\Phi_i^{-1}(u_i)$, $u_i \in [0, 1]$, Corollary 1 (Equation 3.13) implies that the normal copula

$$C_n(\mathbf{u} \mid \mathbf{P}_n) = \Phi_3(\Phi_1^{-1}(u_1), \Phi_2^{-1}(u_2), \Phi_3^{-1}(u_3) \mid \mathbf{P}_n). \quad (3.33)$$

Thus, the normal copula equals

$$\int_{-\infty}^{\Phi^{-1}(u_1)} \int_{-\infty}^{\Phi^{-1}(u_2)} \int_{-\infty}^{\Phi^{-1}(u_3)} |(2\pi\mathbf{P}_n)|^{-1/2} \exp\left(-\frac{1}{2}\mathbf{x}^T \mathbf{P}_n \mathbf{x}\right) d\mathbf{x}, \quad (3.34)$$

where $\mathbf{x} \in \mathbb{R}^3$. By Sklar’s theorem, coupling the normal copula $C_n(\mathbf{u} \mid \mathbf{P}_n)$ with three normal marginals $F_i(x_i)$ will result in the multivariate normal distribution function $C(F_1(x_1), F_2(x_2), F_3(x_3) \mid \mathbf{P}_n)$, but substituting some other univariate marginals $G_i(x_i)$ into the normal copula will result in a novel probability distribution $C_n(G_1(x_1), G_2(x_2), G_3(x_3) \mid \mathbf{P}_n)$.

Derivation of a t copula is also a consequence of applying Corollary 1. Given a 3-dimensional t cumulative distribution function $T_3(\cdot)$ and a t quantile function $t_i^{-1}(u_i)$, the t copula

$$C_t(\mathbf{u} \mid \mathbf{P}_t, \omega) = T_3(t_1^{-1}(u_1 \mid \omega), t_2^{-1}(u_2 \mid \omega), t_3^{-1}(u_3 \mid \omega) \mid \mathbf{P}_t), \quad (3.35)$$

where \mathbf{P}_t is a positive definite correlation matrix like \mathbf{P}_n , and $\omega > 0$ is a degrees of freedom parameter. Definitions of $T_3(\cdot)$ and $t_i^{-1}(\cdot)$ imply that the t copula equals

$$\begin{aligned} \int_{-\infty}^{t^{-1}(u_1|\omega)} \int_{-\infty}^{t^{-1}(u_2|\omega)} \int_{-\infty}^{t^{-1}(u_3|\omega)} & \mathcal{G}([\omega + 3]/2) / \left[\mathcal{G}(2\omega) \sqrt{|\pi\omega\mathbf{P}_t|} \right] \\ & \times \left(1 + \mathbf{x}^T \mathbf{P}_t^{-1/2} \mathbf{x} / \omega \right)^{(\omega+3)/2} d\mathbf{x}, \end{aligned} \quad (3.36)$$

where $\mathcal{G}(\cdot)$ is a univariate gamma function and $\mathbf{x} \in \mathbb{R}^3$. Unlike the normal copula, the t copula has a degrees of freedom parameter ω that results in heavier tails. The parameter ω introduces additional flexibility that controls the probability of extreme values in the tails, resulting in positive tail dependence, and combinations of parameters orthogonal to the correlation structure.

As an example of orthogonality, supposing that correlation between two parameters is positive, it is unlikely that we might observe a pair of values where one point has high magnitude and the second point has low magnitude. However, the occurrence of such rare events is predicted for finite values of ω with considerably higher probability than for the normal copula. When $\omega \rightarrow \infty$, the t copula converges to the normal copula (Joe, 1997). Thus, a model of parameter variation based on the t copula is a further generalization

of the normal copula that allows modeling correlations between processing components that sometimes take unusual combinations or extreme values in the tails.

To complete definitions of two models of parameters, Sklar's theorem requires that we specify marginal distributions for each parameter. Because I am modeling parameters in a descriptive manner and variation of parameters is unobservable, I limited my search to univariate distributions that exhibit flexible shapes to enable good approximations of the true shapes of parameter densities. In contrast, if I were following a theoretical approach, I would want to produce marginal variation that is derived theoretically and disallow all other distributional shapes.

For the drift rate δ , consistent with the Ratcliff model, I will assume a normal distribution. I will replace Ratcliff's continuous uniform assumption for the non-decision parameter τ^{nd} with a truncated normal distribution, and place a beta distribution on the bias parameter β . The univariate marginals for the three parameters are therefore

$$\begin{aligned}\delta &\sim \mathcal{N}(\nu, \eta^2), \\ \beta &\sim \mathcal{B}(\lambda, \gamma), \text{ and} \\ \tau^{nd} &\sim \mathcal{TN}(\chi, \phi^2, a, b),\end{aligned}\tag{3.37}$$

where $\lambda > 0$ and $\gamma > 0$ are shape parameters, $\chi \geq a$ is a location parameter, $\phi > 0$ is a scale parameter, a and b are truncation points such that $0 \leq a < b < \infty$.

Now we can combine specified copulas and univariate marginals. Apply-

ing Sklar's theorem leads to $i = 1, 2$ new multivariate distributions of the form

$$F_i(\delta, \beta, \tau^{nd}) = C_i(F_1(\delta), F_2(\beta), F_3(\tau^{nd})), \quad (3.38)$$

from which, by Corollary 2, we can obtain probability densities by

$$f_i(\delta, \beta, \tau^{nd}) = c_i(F_1(\delta), F_2(\beta), F_3(\tau^{nd}))f_1(\delta), f_2(\beta), f_3(\tau^{nd}). \quad (3.39)$$

The copula density and marginal densities of parameters have known closed-form expressions (Nelsen, 2007; Casella and Berger, 2002; Johnson et al., 1994, 1995). Direct application of Sklar's theorem ensures that both models of parameter variability are valid distributions or densities.

This completes the generalization of the Ratcliff diffusion model of simple decisions. We can now proceed to study implications of the proposed models.

Chapter 4

Proposed studies

With models of behavior defined, I can now outline two studies designed to answer the remaining questions. Recall that I proposed to compare predictions of the models and analyze the nature of the correlation structure consistent with the benchmark data. In Study A, I plan to explore effects of parameter correlations on models' predictions under an experimental design similar to Ratcliff and Rouder (1998). The objective is to compare predictions of mean sample paths, and commonly found features of response and response times that follow from their joint densities.

Study B will apply the proposed models to the benchmark dataset. Posterior densities of parameters and predictions will be obtained with MCMC. The goals are to infer the directions and plausible values of correlation parameters, compare the quality of predictions of fitted models for important features of behavioral data and determine the effect of adding a correlation structure on a model complexity-adjusted measure of overall predictive ac-

curacy.

4.1 Study A: Comparison of Models' Predictions

Study A will examine predictions of the standard Ratcliff diffusion models and its generalizations proposed above. I will first outline exactly what predictions I propose to use for model comparison. It turns out that all these predictions require solutions of analytically intractable integration problems, so along with the predictions I will describe methods of approximation appropriate for each problem. Then, I will describe the experimental design and values of free parameters under which predictions will be calculated. Finally, I will describe graphical and quantitative means of comparing the three models and discuss limitations of this study.

4.1.1 Mean Sample Paths

One prediction of the three decision models that can provide insight into effects of correlations among parameters of a Ratcliff diffusion model include mean sample paths. The mean sample path $E[x(t)]$ is an unconditional mean of a stochastic process at each time point t . For the Wiener process with no boundaries, it is available in analytical form and is equivalent to $\delta t + \zeta$, where δ is the drift rate the slope and ζ is the starting point. However, for a Wiener process with absorbing boundaries, as required for the decision process model,

the dynamics change in a non-trivial way and the mean sample path is not available in explicit form. The calculation is further complicated by variable parameters.

However, we can still obtain predictions by realizing that the mean sample path is an integral problem at each time point. Consider the value of the sample path of a bounded Wiener process with fixed parameters at some time t . Then, the mean sample path

$$\mathbb{E}[x(t)] = \int_{\mathcal{S}} x(t) f(x(t)) dx(t), \quad (4.1)$$

where $\mathcal{S} = [0, \alpha]$ is the bounded state space and $f(x(t))$ is the marginal density of $x(t)$. Because $f(x(t))$ is unavailable, we cannot use numerical integration. However, the problem could be solved if we could sample from $f(x(t))$ and then use a Monte Carlo estimator to approximate the sample path at time t .

The Monte Carlo estimator appropriate to this problem is a consequence of a law of large numbers for independent and identically distributed random variables (Casella and Berger, 2002). For a random sample $x_1(t), x_2(t), \dots, x_N(t) \sim f(x(t))$, with expected value $\mathbb{E}[x(t)] < \infty$,

$$\frac{1}{N} \sum_{i=1}^N x_i(t) \cong \mathbb{E}[x(t)], \quad (4.2)$$

when N is large.

To obtain a Monte Carlo estimate of the mean sample path we need samples from the marginal density over a grid of time points. One way to obtain such samples is to simulate the Wiener process until it hits an

absorbing boundary. Each simulated path provides a draw from the marginal density at each time point of the grid. Because sample paths are generated independently, values of the process at time t make an independent sample. Hence, assumptions for the Monte Carlo estimator hold for each time point. The issue, however, is how to average the sample paths given that they will have different lengths.

One way to resolve the length issue is to align paths either on the initial time or the absorption time, and average paths along the length of the shortest path. This way one can estimate either the initial or final part of the mean sample path. I propose aligning at the absorption time because that is where the dynamics are different from the unbounded process and effects of parameter correlations will be most visible as they build up over the time course of a decision.

The above considerations were made under the assumption of fixed parameters, but variable parameters pose no difficulty. The Ratcliff diffusion model and its generalizations assign some multivariate density over the parameters. To obtain a mean sample path requires averaging a mean sample path for fixed parameters with respect to the parameter density. Because it is just another layer of integration, the Monte Carlo method will automatically calculate it if we simulate it in a hierarchical manner.

In a hierarchical simulation, we need to incorporate variability in parameters. For the decision model $i = 1, 2, 3$, a hierarchical procedure to obtain estimates of the mean sample paths will involve

1. Simulating N draws of $(\nu, \zeta) \sim f_i(\nu, \zeta \mid \nu, \eta^2, \lambda, \gamma)$,
2. Simulating N sample paths $x(t \mid \nu, \zeta, \alpha)$, and
3. Aligning and averaging sample paths from step 2 along the shortest path.

The procedure for estimating mean sample paths is computationally simple and fast. Simulation in Step 1 can be obtained with the `base` and `copula` packages in the R statistical environment (Hofert et al., 2013; R Core Team, 2014). For Step 2, I will use a custom R program for a random walk approximation of the Wiener process (Tuerlinckx et al., 2001).

To conclude, understanding of how correlated parameters affect the time course of a decision process can be revealed by comparing mean sample paths of the standard to generalized Ratcliff diffusion models. The calculation of $E[x(t)]$ is very difficult, but we can obtain it by simulating the bounded Wiener process with variable parameters and then averaging sample paths. From examining mean sample paths we can move to studying effects of correlated parameters on response times and responses.

4.1.2 Response times and Responses

While studying mean sample paths under correlated parameters will suggest how the decision process changes for different patterns of correlation, examining the joint density of response times and responses will show how behavior changes. Recall that the joint density is a mixture model averaged

with respect to variability in parameters, such that for model $i = 1, 2, 3$

$$f_i(t^{rt}, r \mid \boldsymbol{\theta}_i) = \int_A f(t^{rt}, r \mid \alpha, \delta, \zeta) f_i(\delta, \zeta \mid \boldsymbol{\theta}_i) dA, \quad (4.3)$$

where $\boldsymbol{\theta}_i$ will reflect the correlation structure and other parameters defined as above.

For a fixed correlation structure, the joint density of behavior contains all the predictions of interest. We could look at its shape, which consists of two density components, one for each response. Or we could examine summary features like the vector of mean of response times and responses, and their covariance matrix. However, a more popular approach is to examine response time density function conditioned on choice (Vickers, 1979; Luce, 1986), both of which follow from the joint density of t^{rt} and r . I will adopt the standard approach because it matches the often observed performance asymmetry between different responses and provides a systematic way of examining effects of a correlation structure.

To determine effects of a correlation structure I propose to examine shapes of conditional densities and their various summary features. We can obtain conditional densities by normalizing the joint density with the probability of response. For response r and model i , conditional density of response time is

$$f_i(t^{rt} \mid r, \boldsymbol{\theta}_i) = f_i(t^{rt}, r \mid \boldsymbol{\theta}_i) / P\{r = c\}, \quad (4.4)$$

where $c \in \{0, 1\}$ is a choice. One feature of conditional densities worth examining is their overall shape. Observed shapes are positively skewed and change in characteristic ways across experimental conditions. The correlation

structure should have effects on the shape and provide a qualitative way of comparing the models.

In contrast to qualitative comparisons, we can also calculate summary features of each conditional density that quantitatively characterize their shapes. To provide an informative account of effects of correlated parameters it is necessary to capture location, scale, shape and the corresponding response probabilities (normalizing constants) of each conditional density of response time. For location I will use the expected value as a measure of location, $E_i[t^{rt} \mid r] = \mu_{r,i}$, and standard deviation as a measure of scale, $\sqrt{E_i[(t^{rt} - \mu_{r,i})^2 \mid r]}$. The shape is a complex concept with many components, but here I propose to use skew,

$$\text{Skew}_i[t^{rt} \mid r] = \frac{E_i[(t^{rt} - \mu_{r,i})^3 \mid r]}{E_i[(t^{rt} - \mu_{r,i})^2 \mid r]^{3/2}}, \quad (4.5)$$

and excess kurtosis,

$$\text{Kurtosis}_i[t^{rt} \mid r] = \frac{E_i[(t^{rt} - \mu_{r,i})^4 \mid r]}{E_i[(t^{rt} - \mu_{r,i})^2 \mid r]^2} - 3. \quad (4.6)$$

Data show consistently positive skew of response times, which can be measured by skew, and excess kurtosis can capture increase in the amount of extreme response times relative to what can be expected under a normal density. Last, the probability of hitting the upper boundary

$$P_i\{r = 1\} = \int_{\delta \times \beta} \left[\exp\left(\frac{-2\delta\alpha\beta}{\sigma^2}\right) - 1 \right] / \left[\exp\left(\frac{-2\delta\alpha}{\sigma^2}\right) - 1 \right] f_i(\delta, \beta \mid \boldsymbol{\theta}_i) d\delta d\beta, \quad (4.7)$$

where the probability of hitting the lower boundary is $P_i\{r = 0\} = 1 - P_i\{r = 1\}$.

Notice that computing the conditional densities and calculating all the summary features again requires integration. Similar to the problem of finding mean sample paths, the integrals are intractable, but in this case the probability density is known and computable (Tuerlinckx, 2004; Navarro and Fuss, 2009). Instead of using a Monte Carlo estimate, we can use deterministic approximation methods based on quadrature. Similar to the Monte Carlo estimator, quadrature methods approximate an integral with a weighted average, such that

$$\int_{\mathcal{X}} g(x)dx \approx \sum_{i=1}^N w_i g(x_i), \quad (4.8)$$

but weights w_i can be other than $1/n$ and nodes x_i are set deterministically.

To obtain behavioral predictions, I propose using a globally adaptive integration algorithm (Genz and Malik, 1980; Berntsen et al., 1991). It is highly efficient for 2- and 3-dimensional integrals, and precision can be guaranteed deterministically. Optimized code has been implemented in C and can be called using R package cubature (Johnson and Narasimhan, 2013).

With predictions worth exploring outlined, we now need to identify a task and experimental conditions under which predictions for decision process and behavior can be compared. In the next section, I outline how to obtain predictions for a brightness discrimination task used by Ratcliff and Rouder (1998).

4.1.3 Parameter Settings

By now, predictions of interest and methods for their calculations have been specified. Next, I lay out experimental conditions under which predictions will be obtained. I follow the principle that predictions should be obtained for a particular task under realistic experimental conditions. Because models will later be fit to the benchmark dataset it is most natural to expore predictions for the underlying brightness discrimination task.

Recall that the Ratcliff and Rouder (1998) benchmark dataset was generated by a two-choice free response perceptual task under 33 brightness levels and speed-accuracy instructions. They had 3 participants, but they showed similar behavior and parameters, so I will concentrate on a single individual. To calculate predictions we need parameter values that reflect experimental manipulations experienced by the individual. One approach could be to estimate parameters representing effects of manipulations from a range of plausible values that were recovered over many experiments (Matzke and Wagenmakers, 2009). In contrast, it is a benchmark dataset, there are parameter estimates for the Ratcliff diffusion model available, both Bayesian and non-Bayesian (Ratcliff and Rouder, 1998; Vandekerckhove et al., 2011). I propose using the means of posterior densities from Vandekerckhove et al. (2011) to obtain predictions.

Bayesian modeling done by Vandekerckhove et al. (2011) made certain assumptions about how experimental manipulations influence parameters. In their best-fitting model, the authors assumed that brightness level affects

only mean drift rate and speed-accuracy instructions affect drift rate, boundaries and starting position, with the other parameters being constant across the experimental sessions. I will calculate predictions under these assumptions, too.

For the drift rates, Vandekerckhove et al. (2011) provided a descriptive front-end model of how the rate of evidence accumulation arises from brightness levels under particular instructions. The formal relation is described by the Weibull function. For instruction $i = 1, 2$ and brightness level $j = 1, 2, \dots, 11$, I assume mean drift rate regresses on brightness level through a sigmoidal-shaped Weibull function

$$\nu_{i,j} = \nu_i^{lo} + (\nu_i^{hi} - \nu_i^{lo}) \left(1 - \exp \left([-p_j / \nu_i^{sc}] \nu_i^{sh} \right) \right), \quad (4.9)$$

where $\nu_{i,j}$ is the drift rate, p_j is the brightness proportion, ν_i^{lo} is the lower asymptote, ν_i^{up} is the upper asymptote, ν_i^{sc} is the scale parameter and ν_i^{sh} is the shape parameter. The estimates of Weibull parameters taken from Vandekerckhove et al. (2011) are in Table 4.1.

Instruction	ν_{lo}	ν_{hi}	ν_{sh}	ν_{sc}
Accuracy	-0.352	0.329	4.413	0.526
Speed	-0.565	0.511	5.227	0.521

Table 4.1: Weibull function parameters

The Weibull function provides a parsimonious way of connecting 66 experimental conditions to drift rates via 8 parameters. However, for obtaining predictions I do not need all 66 conditions. Based on previous fits to

the benchmark dataset (Ratcliff and Rouder, 1998; Vandekerckhove et al., 2011), and more commonly observed in psychophysical investigations (Vickers, 1979; Ratcliff, 2014), drifts and discriminability performance level off for extremely bright or dark stimuli. The drift rates change sigmoidally as a function of brightness level and are most informative for proportions from 0.2 to 0.8. Hence, I can use points in that range that were also used in the benchmark experiment to set drift rates. More specifically, I propose to use 11 white-to-black pixel proportions, $\{0.188, 0.250, 0.313, 0.375, 0.438, 0.500, 0.563, 0.625, 0.688, 0.750, 0.813\}$, which will give a sufficient number of points to capture the shape of predicted curves characterizing participant performance.

The other non-correlation parameters are listed in Table 4.2. The values for thresholds α_i and drift rate standard deviation η are taken directly from Vandekerckhove et al. (2011). Mean bias λ_i and standard deviation of bias γ are transformations of end point estimates of a uniform random variable because I assumed a beta density whereas Vandekerckhove et al. (2011) assumed a continuous uniform density. Similarly, Vandekerckhove et al. (2011) assigned a normal density to the non-decision parameter whereas I assigned a normal truncated at 0, so their estimates of mean and standard deviation now correspond to my location and scale parameters of the truncated normal.

Instruction	α	η	λ	γ	χ	ϕ
Accuracy	0.221	0.127	0.464	0.065	0.279	0.041
Speed	0.050	0.127	0.464	0.008	0.279	0.041

Table 4.2: Non-correlation parameters

The remaining parameters are the degrees of freedom for the t copula and correlation parameters that characterize generalized diffusion models. I will fix the degrees of freedom ω to 5, which gives a t copula having all its moments and being maximally different from the normal copula. The ω parameter will be assumed to be the same throughout the experiment.

Last, I propose to examine three correlation structures with low, medium and high levels of correlation between parameters, which will also stay the same throughout the experiment. The patterns are given in Table 4.3. The structures reflect a few of many possible correlation patterns that could exist under typical experimental conditions. The selected patterns can be interpreted to represent different processing situations. For example, based on the pattern of signs, sets 1 - 3 can be interpreted as a tendency of decision bias β and evidence δ to agree on the best response while the non-decision τ^{nd} reflects motor bias towards the other response. Thinking through the behavioral consequences of other correlation structures, similar interpretations can be proposed, but these are only suggestive and were used to select the three proposed correlation patterns out of the infinitely many possibilities without any theoretical commitment.

Set	$\rho_{\delta,\beta}$	$\rho_{\delta,\tau^{nd}}$	$\rho_{\beta,\tau^{nd}}$
1	0.15	-0.15	-0.15
2	0.50	-0.50	-0.50
3	0.85	-0.85	-0.85
4	-0.15	0.15	0.15
5	-0.50	0.50	0.50
6	-0.85	0.85	0.85
7	-0.15	-0.15	0.15
8	-0.50	-0.50	0.50
9	-0.85	-0.85	0.85

Table 4.3: Correlation parameters

This combination of proposed parameters should lead to interpretable comparisons between the three models. All calculations will be done under the same parameter values for all the non-correlation parameters, and correlation parameters will be the same for the two copula-based models. The only differences between the models are the copula assumptions. The Ratcliff diffusion model has the independence copula and generalizations have the normal and t copulas.

Moreover, the copula-based models will be matched in their Kendall coefficients for each correlation pattern because, as mentioned in the copula section, they both belong to a class of elliptical distributions whose Kendall's coefficient is equal to $2 \arcsin(\rho)/\pi$. However, recall that the t copula is

nevertheless distinct by having symmetric tail dependence (the asymptotic probability that two random variables are both extreme). For the t copula, the tail dependence is equal to

$$2t_{\omega}(-\sqrt{\omega+1}\sqrt{1-\rho}/\sqrt{1+\rho}), \quad (4.10)$$

where t_{ω} is a t cumulative distribution function, ω is its degrees of freedom and ρ is the correlation parameter. The tail dependence is determined by both parameters ω and ρ , which gives a more concrete justification for choosing $\omega = 5$ since it ensures that t copula is distinct from a normal copula in its tail dependence.

In conclusion, the proposed set of parameters is intended to enable unambiguous conclusions about how the three models differ in their predictions based solely on their parameter dependence assumptions. In the next section I will describe several means of organizing predictions to facilitate comparison and establish differences between models.

4.1.4 Results

How do the three decision models differ in their predictions? We can answer this question in detail graphically and quantitatively. After calculating mean sample paths, conditional densities of response times and their summary features, I will organize them graphically and provide calculations of prediction differences where models diverge the most.

Trajectories of mean sample paths naturally form a time series. To understand how correlation between ζ and δ affects shapes of sample paths I

will overlay Monte Carlo estimates for the three models for each combination of brightness and instructions. The combination of 22 plots will reveal qualitatively how the time course of the decision process is shaped by each of 9 correlation structures.

The conditional densities of response times can also be graphically compared to explore effects on the overall shape of those densities. Juxtaposing conditional densities using violin plots will reveal asymmetries between the two choices. The full pattern of variation in the density shape will be examined across combinations of experimental condition and correlation structure.

The shapes of conditional densities can be compared quantitatively using various summary features. For each density, there will be a mean, standard deviation, skew, kurtosis and probability of the corresponding response. I will examine how the summary features vary as a function of brightness level and instruction for each density. From the same prediction data I will also examine simultaneous changes between densities using latency-probability plots that relate mean conditional reaction times with probability correct. In addition to the plots, since quantities for all the features will be known, I can also provide precise discrepancies between models for particular conditions.

Overall, the full pattern of graphical and quantitative comparisons of the three models will reveal effects of parameter correlation in detail. The detailed understanding of prediction deviations will enable us to judge which features of data to which we should pay attention when evaluating models against the benchmark dataset.

4.1.5 Discussion

Results of model predictions will provide a rich body of data to discuss the assumption of correlated parameters. Whatever these conclusions may be they will be limited in their generality by certain choices underlying this study. The two major sources of limitation are the modeling assumptions and the response time task assumptions. To examine effects of correlated parameters I propose to use the Ratcliff diffusion model, but other well-established sequential sampling models could be used, too (Usher and McClelland, 2001; Brown and Heathcote, 2008). Certain model features such as noisy evidence or perfect integration without inhibition, may interact with correlated parameters in a qualitatively different manner from when these assumptions do not hold. Due to the sensitivity of decision models to architectural assumptions (Luce, 1986; Ratcliff and Smith, 2004), conclusions about predictions obtained in this thesis will not generally be transferable to other models.

In addition to the accumulation process, using elliptical copulas to represent correlation structure is also a limitation. In elliptical copulas, while each pair of dimensions can have different signs and magnitudes of correlation, they all have the same linear dependence structure. A more flexible approach is to model the dependence between each pair of variables by a different bivariate copula and tie them together into a novel multivariate copula using a tree-like factorization of a multivariate distribution (Czado, 2010). If there are some data, such as trial-level Bayes estimates of parameters, that reveal the dependence structure, a custom copula should make better

predictions.

Another major limitation comes from obtaining model predictions under the experimental design and manipulations of the Ratcliff and Rouder (1998) study. As discussed in Section 1.2 on page 21, other response time procedures, like the response signal procedure, could have been used or manipulations of rewards might be added or the number of choices increased. The proposed predictions will be limited to two-choice free response procedure with manipulations of stimulus quality and speed-accuracy instruction.

4.2 Study B: Benchmark Dataset Analyses

The purpose of Study B is to evaluate the proposed models against the benchmark behavioral data and estimate the correlation structure of their parameters. The outline of the study is as follows. First, I will describe an outlier pre-processing procedure to deal with extreme response times. Next, I will define three hierarchical Bayesian models that will be fit to the pre-processed data. A variant of a Metropolis-Hastings algorithm will be defined and convergence diagnostics discussed. Finally, I will describe statical analyses that can provide an answer to the study questions and discuss some limitations of the study.

4.2.1 Outlier Pre-processing

An important feature of any dataset involving response times is presence of contaminants. Contaminants are observations arising from processes unre-

lated to research questions. Some of the recognized types of contaminants are fast guesses, slow guesses and delayed responses (Ratcliff, 1993; Ratcliff and Tuerlinckx, 2002; Vandekerckhove and Tuerlinckx, 2007, 2008; Craigmile et al., 2010; Vandekerckhove et al., 2011). Contaminants could arise due to non-compliance with experimental instructions or occasional loss of attention, among other processes.

To obtain accurate inferences about the process of interest, contaminants need to be either excluded or modeled. The proposed models are complex and initial testing may be more successful without including additional structure required to capture contaminants (Ratcliff and Tuerlinckx, 2002; Craigmile et al., 2010). Therefore, I will use a set of common procedures to clean the response time data before fitting the models (Ratcliff, 1993).

Because contaminants mixed with regular response times cannot be identified, the most that is possible to do is to remove outliers. An outlier is defined as an extreme observation relative to the mean of the sample. There are fast and slow outliers that appear in a typical dataset. The fast outliers are usually to be understood as fast guesses, thus a useful heuristic to detect them is chance performance (Ratcliff and Tuerlinckx, 2002).

One way to use chance performance to remove fast outliers is to bin data into non-overlapping, continuous intervals. Starting at the lowest multiple of the bin size that includes the fastest response times, we can scan the ordered sequence of observations one interval at a time. For each interval, we calculate the accuracy rate. If accuracy is at chance, then responses are fast guesses and they should be excluded. I will define “at chance” as

an interval between 0.50 and 0.53, taking into account sampling error and the fact that the probability of exactly chance performance is zero. When accuracy rises above the chance interval, then responses in the bin will be assumed to be due to the process of interest and the lower bound of the bin forms the cut-off for excluding fast responses.

For the benchmark dataset, I propose using a window of 50 ms. The individuals and conditions vary a lot, so the method will be more accurate if it is applied to all the participant-by-condition combinations as opposed to applying it to the pooled data. A preliminary application of the method resulted in a cut-off value of 200 ms for all three participants. The total number of observations excluded with this cutoff is 123.

On the other hand, dealing with extremely long reaction times is different because it includes both slow guesses and delayed responses. In this case, the accuracy signature is not appropriate because delayed responses can be highly accurate because they are generated by the decision process. A practical method to deal with slow outliers is to use a multiple of standard deviations from the mean as a cut-off.

To remove the slow outliers I propose using three standard deviations from the mean as a cut-off. As for the fast outliers, I will apply the procedure to all participant-by-condition combinations. Preliminary application of the cut-offs led to removal of 429 observations. Overall, applying these proposed procedures would bring the total number of excluded observations to 2% of the full dataset.

4.2.2 Bayesian Models

With the data to be modeled established, the remaining task is to describe models to be fitted. The models of performance data I plan to evaluate are hierarchical versions of the independent copula diffusion model, the normal copula diffusion model and the t copula diffusion model. Multiple variations of these models could be constructed by making different assumptions about how experimental manipulations, stimulus quality and speed-accuracy instructions influence cognitive processing. I will rely on results from previous modelling and an EEG study to hypothesize how experimental manipulations affect processing. Finally, while a hierarchy can be defined for all parameters of the model, I will restrict it to parameters of primary interest: these parameters that quantify experimental manipulations and the correlation parameters.

I will use the following variables to denote features of the experimental design. The dataset contains $l = 1, 2, 3$ participants. Each participant received a block-wise instruction $k = 0, 1$, where 0 is speed and 1 is accuracy, combined with a trial-wise brightness level $j = 0, 1, \dots, 32$. Therefore, for each condition, a participant would have $N_{j,k,l}$ observations. The observations come from a model $M = 1, 2, 3$, referring to independent, normal and t diffusion model, respectively.

Recall that the joint density of response times and responses is a mixture of the Wiener first-passage density and a copula-based density. The joint density for model M of a single observation $(t_{i,j,k,l}^{rt}, r_{i,j,k,l})$ conditioned on

instruction k and brightness proportion p_j is

$$f_M(t_{i,j,k,l}^{rt}, r_{i,j,k,l} \mid p_j, \alpha_{k,l}, \nu_{j,k,l}, \eta_l, \lambda_{k,l}, \gamma_{k,l}, \chi_{k,l}, \phi_{k,l}, \rho_l^{\delta, \tau^{er}}, \rho_l^{\beta, \delta}, \rho_l^{\beta, \tau^{er}}). \quad (4.11)$$

The indices of parameters reflect individual differences across all parameters. The instructions are assumed to influence thresholds, drift rates and bias, as with previous successful model fits of the benchmark data and similar experiments (Ratcliff and Rouder, 1998; Vandekerckhove and Tuerlinckx, 2007; Ratcliff and McKoon, 2008; Wagenmakers, 2009; Vandekerckhove et al., 2011), and also influence nondecision times as indicated by an increase in the lateralized readiness potential under accuracy instructions (Rinkenauer et al., 2004). The brightness level only influences drift rates (Voss et al., 2004; Ratcliff and McKoon, 2008; Vandekerckhove et al., 2011).

Similar to Study A, I will model the effect of brightness proportion and instruction on drift rates with a regression function. I assume that mean drift rate regresses on brightness proportion through a sigmoidal-shaped Weibull function

$$\nu_{j,k,l} = \nu_{k,l}^{lo} + (\nu_{k,l}^{hi} - \nu_{k,l}^{lo}) \left(1 - \exp \left([-p_j / \nu_{k,l}^{sc}]^{\nu_{k,l}^{sh}} \right) \right), \quad (4.12)$$

where $\nu_{j,k,l}$ is the drift rate, s_j is the brightness level, $\nu_{k,l}^{lo}$ is the lower asymptote, $\nu_{k,l}^{up}$ is the upper asymptote, $\nu_{k,l}^{sc}$ is the scale parameter and $\nu_{k,l}^{sh}$ is the shape parameter. The parameter indices express individual differences and effect of speed/accuracy manipulation on all the parameters (Vandekerckhove et al., 2011). This completes specification of model M's data density.

The next part of the Bayesian models I specify are the priors for all the parameters. All parameters except for the correlation parameters and the de-

degrees of freedom are shared by the three models. I will assign the same priors across models to all the shared parameters to maximize similarity. Experience with fitting the Ratcliff model across many tasks and conditions suggests a range of plausible values that can be incorporated through restricting support of priors. I will use recovered parameters from 28 studies reported by Matzke and Wagenmakers (2009), but double the reported maximum values to take into account the possibility of yet unobserved rare parameter combinations or imperfect representativeness of tested participants. The imposed constraints should improve extraction of parameter information from data by reducing posterior variability (Gelman et al., 2013a).

To impose flexible support constraints while allowing for adaptable distribution of probability mass over the domain, I will use truncated normal priors for the main parameters of interest. The priors for the commonly shared main parameters of interest are

$$\begin{aligned}
\alpha_{k,l} &\sim \mathcal{TN}(\mu_k^\alpha, \sigma_k^\alpha, 0.001, 0.786), \\
\nu_l^{lo} &\sim \mathcal{TN}(\mu^{\nu^{lo}}, \sigma^{\nu^{lo}}, -1.172, 0), \\
\nu_l^{up} &\sim \mathcal{TN}(\mu^{\nu^{up}}, \sigma^{\nu^{up}}, 0, 1.172), \\
\nu_l^{sh} &\sim \mathcal{TN}(\mu^{\nu^{sh}}, \sigma^{\nu^{sh}}, 0, 50), \\
\nu_l^{sc} &\sim \mathcal{TN}(\mu^{\nu^{sc}}, \sigma^{\nu^{sc}}, 0, 5), \\
\lambda_{k,l} &\sim \mathcal{TN}(\mu_k^\lambda, \sigma_k^\lambda, 0, 1), \text{ and} \\
\chi_{k,l} &\sim \mathcal{TN}(\mu_k^\chi, \sigma_k^\chi, 0, 1.884).
\end{aligned} \tag{4.13}$$

The normal and t diffusion model add a correlation structure to the inde-

pendent model, so the correlation parameters for $M = 2, 3$ will have priors such that

$$\begin{aligned}\rho_l^{\delta, \tau^{er}} &\sim \mathcal{TN}(\mu^{\rho^{\delta, \tau^{er}}}, \sigma^{\rho^{\delta, \tau^{er}}}, -1, 1), \\ \rho_l^{\beta, \delta} &\sim \mathcal{TN}(\mu^{\rho^{\beta, \delta}}, \sigma^{\rho^{\beta, \delta}}, -1, 1), \text{ and} \\ \rho_l^{\beta, \tau^{er}} &\sim \mathcal{TN}(\mu^{\rho^{\beta, \tau^{er}}}, \sigma^{\rho^{\beta, \tau^{er}}}, -1, 1),\end{aligned}\tag{4.14}$$

with support from -1 to 1 reflecting ignorance about correlations. Last, the t diffusion model has the additional degrees of freedom parameter, so when $M = 3$ the density over parameters will be augmented with

$$\omega_l \sim \mathcal{TN}(\mu^\omega, \sigma^\omega, 0, \infty),\tag{4.15}$$

where support from 0 to ∞ also reflects ignorance about plausible values of the parameter.

Up to now, I have specified the parts of the overall prior probability density that are composed of the main parameters of interest. The remaining parameters describe across-trial variability. I will adopt the simplifying assumption that experimental manipulations have no effect on these parameters and that there are no similarities among participants with respect to them, which simplifies the hierarchical structure. In selecting these priors, I maintain the approach of restricting support to plausible values, but select uniform priors to simplify computation. Then, the priors of across-trial variability parameters are

$$\begin{aligned}
\eta_l &\sim \mathcal{U}(0, 0.658), \\
\gamma_l &\sim \mathcal{U}(0, 0.860), \text{ and} \\
\phi_l &\sim \mathcal{U}(0, 1.260).
\end{aligned}
\tag{4.16}$$

To finish defining the target Bayesian models, we need priors on second level parameters. My approach is to assign restricted uniform densities to the location parameters of the truncated normal priors and assign mildly informed half-Cauchy densities to scale parameters (Polson et al., 2012; Gelman et al., 2013a). The functional forms of second-level priors fit naturally with the rest of the model because location parameters have domains equivalent to the support of the priors, and scale parameters can take any positive real value. This approach allows all possible parameter combinations to be considered with minimal a priori commitment to any region of the parameter space, reflecting the belief that parameters are not the same across experimental tasks.

For the three models, shared location parameters will be assumed to

follow

$$\begin{aligned}
\mu^\alpha &\sim \mathcal{U}(0.001, 0.786), \\
\mu^{\nu^{lo}} &\sim \mathcal{U}(-1.172, 0), \\
\mu^{\nu^{up}} &\sim \mathcal{U}(0, -1.172), \\
\mu^{\nu^{sh}} &\sim \mathcal{U}(0, 50), \\
\mu^{\nu^{sc}} &\sim \mathcal{U}(0, 5), \\
\mu^\lambda &\sim \mathcal{U}(0, 1), \text{ and} \\
\mu^\chi &\sim \mathcal{U}(0, 1.884),
\end{aligned} \tag{4.17}$$

and shared scale parameters will follow

$$\begin{aligned}
\sigma^\alpha &\sim \mathcal{HC}(25), \\
\sigma^{\nu^{lo}} &\sim \mathcal{HC}(25), \\
\sigma^{\nu^{up}} &\sim \mathcal{HC}(25), \\
\sigma^{\nu^{sh}} &\sim \mathcal{HC}(25), \\
\sigma^{\nu^{sc}} &\sim \mathcal{HC}(25), \\
\sigma^\lambda &\sim \mathcal{HC}(25), \text{ and} \\
\sigma^\chi &\sim \mathcal{HC}(25).
\end{aligned} \tag{4.18}$$

For the normal and t diffusion models, correlation parameters will have location

$$\begin{aligned}
\mu^{\rho^{\delta, \tau^{er}}} &\sim \mathcal{U}(-1, 1), \\
\mu^{\rho^{\beta, \delta}} &\sim \mathcal{U}(-1, 1), \text{ and} \\
\mu^{\rho^{\beta, \tau^{er}}} &\sim \mathcal{U}(-1, 1),
\end{aligned} \tag{4.19}$$

and scale second-level priors

$$\begin{aligned}\sigma^{\rho^{\delta, \tau^{er}}} &\sim \mathcal{HC}(25), \\ \sigma^{\rho^{\beta, \delta}} &\sim \mathcal{HC}(25), \text{ and} \\ \sigma^{\rho^{\beta, \tau^{er}}} &\sim \mathcal{HC}(25).\end{aligned}\tag{4.20}$$

Finally, the degrees of freedom parameter will have location and scale second-level priors

$$\begin{aligned}\mu^{\omega} &\sim \mathcal{HC}(25), \text{ and} \\ \sigma^{\omega} &\sim \mathcal{HC}(25).\end{aligned}\tag{4.21}$$

The proposed hierarchical models will reflect the typical structure of behavioral data. Combining the three participants as members of the same population we can obtain more precise estimates of all parameters, and especially correlation parameters, by semi-pooling all the 8000+ observations of each individual together.

The proposed Bayesian models are complex and require an approximate integration algorithm. In the next subsection, I will describe a variant of the Metropolis-Hastings algorithm with a proposal density motivated by principles of biological evolution.

4.2.3 MCMC Sampler

Selecting a specific Metropolis-Hastings sampler requires specifying a sampling scheme and a proposal density. A sampling scheme defines an iteration

of a sampler in terms of how the dimensions of a Markov chain are updated. It is possible to update all dimensions simultaneously or cycle through them sequentially in a deterministic or random fashion. Once a sampling scheme is chosen, a proposal density will define how we update the position of the Markov chain during an iteration. Both choices are best motivated by considering properties of the model to be fit.

Posterior densities of hierarchical Bayesian models, with a data model derived from psychological theory, tend to have two properties: high-dimensionality and strongly correlated dimensions (Gelman et al., 2013a; Ratcliff and Tuerlinckx, 2002; Turner et al., 2013b). For the Markov chain to move, a proposal density has to generate a candidate vector in high density regions, which becomes harder when dimensions are added and correlations increase. Unless a proposal density is fine-tuned to have scale and covariance parameters that match closely the posterior density, the convergence to an equilibrium density may be very slow. To sample effectively from the proposed models, I propose partitioning the parameter space into blocks of correlated dimensions and proposing candidates based on a rule inspired by evolutionary optimization algorithms.

First, I will introduce the idea of blocked sampling. Recall that standard Metropolis-Hastings updates the p dimensions of a Markov chain simultaneously (also called globally), as shown in Algorithm 1 on page 46. For large p , it may be more difficult to find a proposal that has a high acceptance probability, especially in a correlated space. In this case, replacing simultaneous updating with a sequence of updates over non-overlapping subsets of

dimensions may improve the convergence rate.

As an example, suppose the state space \mathcal{S} can be partitioned into two mutually exclusive blocks $(\mathcal{S}_1, \mathcal{S}_2)$, where $\mathcal{S}_i \in \mathbb{R}^{p_i}$. For block \mathcal{S}_1 , assume there is a Markov chain \mathbf{X}_1 driven by a transition density $p(\mathbf{y}_1 \mid \mathbf{x}_1, \mathbf{x}_2)$, conditioned on the state of \mathbf{X}_2 , to a unique equilibrium density $\pi_{1|2}(\mathbf{x}_1 \mid \mathbf{x}_2)$. Assume the Markov chain over \mathcal{S}_2 has an analogous conditional transitional density $\pi_{2|1}(\mathbf{x}_2 \mid \mathbf{x}_1)$. If the updates of the two chains are executed sequentially, \mathbf{X}_1 followed by \mathbf{X}_2 , then the whole space \mathcal{S} is updated. The full update is represented by the product of the transition densities and turns out to have as its equilibrium density $\pi(\mathbf{x}_1, \mathbf{x}_2)$ (Robert and Casella, 2004; Givens and Hoeting, 2012), that is

$$\iint_{\mathcal{S} \times \mathcal{S}} \pi(\mathbf{y}_1 \mid \mathbf{x}_1, \mathbf{x}_2) \pi(\mathbf{y}_2 \mid \mathbf{y}_1, \mathbf{x}_2) d\mathbf{x}_1 d\mathbf{x}_2 = \pi(\mathbf{y}_1, \mathbf{y}_2). \quad (4.22)$$

This result ensures that breaking up a high-dimensional space into blocks converges to the proper target density and can enhance convergence rate of a Metropolis-Hastings algorithm in a correlated space.

Breaking up the state space into blocks requires specifying a proposal density for each. In principle, blocking is a very flexible sampling scheme because we can design separate proposal densities for each block. However, here I propose using differential evolution candidate generation, based on an evolutionary algorithm for optimization problems of the same name, for all blocks (Storn and Price, 1997; Ter Braak, 2006; Turner et al., 2013b). Evolutionary algorithms perform direct, iterative searches through a space of solutions using principles analogous to biological evolution (Simon, 2013).

Starting with a population of individuals, processes of mutation and cross-over operate on the current individuals to produce a novel individual that competes with the older ones. From the two individuals, the one with better fitness is selected into the next generation. In the final generation the candidate with the best fitness is chosen as a solution.

The structure of biologically inspired optimization algorithms closely maps onto Metropolis-Hastings, but requires some modification. Suppose we want to sample from a target density $\pi(\cdot)$. In the MCMC setting, the solution is a set of points $\mathcal{X} = \{\mathbf{x} : \pi(\mathbf{x}) > 0\}$, as opposed to a single best point, that approximates location, scale and shape of a target density. Application of evolutionary logic suggests that to obtain a good approximation of the target density, we should run K Markov chains that explore the state space in parallel based on a proposal mechanism that perturbs the state of a chain based on the states of other chains. The Differential Evolution MCMC (DE-MCMC) implements this suggested structure as a special case of a Metropolis-Hastings algorithm (Ter Braak, 2006; Turner et al., 2013b).

After blocking the parameter space, DE-MCMC provides a way to approximately sample from the posterior conditional densities. For each block, DE-MCMC starts by independently initiating $K \geq 4$ Markov chains such that initial points are dispersed. Consider a p -dimensional block of parameters that is to be updated. For chain $l = 1, 2, \dots, K$ and iteration $j = 1, 2, \dots, N$, a candidate $\boldsymbol{\theta}^*$ is proposed by perturbing the current state $\boldsymbol{\theta}^{l,j-1}$ with a linear combination of a random “cross-over” process and a random “mutation” process. The cross-over term is a function of two other

chains sampled without replacement and is defined as

$$\nu(\boldsymbol{\theta}^{n,j-1} - \boldsymbol{\theta}^{m,j-1}), \quad (4.23)$$

where $\nu > 0$, $n, m = 1, 2, \dots, K$, $n \neq l$ and $m \neq l$. The mutation term is a p -dimensional random vector $\boldsymbol{\epsilon}$ with support equal to or included in the support of the equilibrium density. Unlike the optimization algorithm, the selection of the proposed candidate is probabilistic and is done using the Metropolis-Hasting acceptance probability. After the acceptance decision, updating of the chain $l+1$ is initiated until DE-MCMC updates all K chains in a fixed order.

An example of DE-MCMC update of a single block is stated as pseudocode in Algorithm 2. The block corresponds to the three correlation parameters $\boldsymbol{\rho} = (\rho_1, \rho_2, \rho_3)$ of proposed generalizations of the Ratcliff model. Block index is suppressed for clarity.

Examining the form of the proposal mechanism suggests why DE-MCMC is useful for highly correlated parameters spaces in psychologically motivated Bayesian models. Recall that to move efficiently in a correlated space requires delicate tuning of the proposal density. The cross-over term of DE-MCMC solves this problem. The underlying assumption is that a population of chains is located in and around regions of high posterior density, which, over time, is guaranteed by Markov process theory. The difference vector $\boldsymbol{\theta}^{n,j-1} - \boldsymbol{\theta}^{m,j-1}$, based on the available population of chains, has magnitude and direction that tends to match the correlation structure of the parameter space. Hence, the proposals are generated in an efficient manner which automatically improves

Algorithm 2 Example DE-MCMC Pseudocode

Require: $\forall l \in \{1, 2, \dots, K\}$, initialize at $\boldsymbol{\rho}^{l,0} \ni \pi(\boldsymbol{\rho}^{l,0} \mid \mathbf{x}) > 0$

for $j = 1, 2, \dots, N$ **do**

for $l = 1, 2, \dots, K$ **do**

 Sample $\boldsymbol{\rho}^{m,j-1}, \boldsymbol{\rho}^{n,j-1} \sim \{\boldsymbol{\rho}^{1,j-1}, \boldsymbol{\rho}^{2,j-1}, \dots, \boldsymbol{\rho}^{K,j-1}\} \setminus \{\boldsymbol{\rho}^{l,j-1}\}$, w.r.

 Sample $v \sim f_v, \boldsymbol{\varepsilon} \sim f_{\boldsymbol{\varepsilon}}, \alpha \sim \mathcal{U}(0, 1)$

 Propose $\boldsymbol{\rho}^* = \boldsymbol{\rho}^{l,j-1} + v(\boldsymbol{\rho}^{m,j-1} - \boldsymbol{\rho}^{n,j-1}) + \boldsymbol{\varepsilon}$

if $\alpha < \pi(\boldsymbol{\rho}^* \mid \mathbf{x}) / \pi(\boldsymbol{\rho}^{l,j-1} \mid \mathbf{x})$ **then**

 Set $\boldsymbol{\rho}^{l,j} = \boldsymbol{\rho}^*$

else

 Set $\boldsymbol{\rho}^{l,j} = \boldsymbol{\rho}^{l,j-1}$

end if

end for

end for

when the placement of chains improves.

Another positive aspect of DE-MCMC is that it has few free parameters that need to be set before sampling can begin. The first is the number of chains K . Recommendations from previous applications suggest that five times the dimensionality of the block works well (Storn and Price, 1997; Ter Braak, 2006; Turner et al., 2013b), so I will adopt this rule. Then, for each chain, there is a scaling parameter v of the cross-over term and parameters $a_\varepsilon, b_\varepsilon$ for the mutation term ε . Similarly, I will adopt recommendations from previous successful applications suggesting random sampling of v from $\mathcal{U}(0.5, 1)$ and to random sampling of ε from $\mathcal{U}(-0.001, 0.001)$.

4.2.4 Convergence Diagnostics

A Metropolis-Hastings algorithm will enable us, in principle, to sample from posterior densities of the proposed models. The Markov chain theory combined with the structure of Metropolis-Hastings will ensure that, given an arbitrary vector of initial values from the parameter space, Markov chains will converge to the equilibrium density. However, the basic theory says nothing about when convergence will take place, information crucial to having accurate inferences.

In a limited number of cases, which are often not useful for real-world applications, there is some formal theory that provides a lower bound on the number of iterations for a given approximation error (Robert and Casella, 2004; Gamerman and Lopes, 2006). More typically, the decision about when

convergence occurred are determined by heuristic methods, often called convergence diagnostics, applied post-hoc to a MCMC sample (Gelman et al., 2013a). The conditions checked by heuristic methods are necessary for convergence, so if they hold it does not mean that convergence holds, but their failure does mean that convergence has not occurred. Working with any difficult sampling problem will require heuristic methods.

We can never be sure about the MCMC output, but we can do graphical and statistical tests that suggest convergence. One popular graphical technique is called a trace plot, which plots the ordered sequence of samples $\{x_1, x_2, \dots, x_N\}$ against the ordering index $i = 1, 2, \dots, N$. Trace plots are usually more convincing when multiple chains starting from overdispersed initial points all converge to the same region and remain overlapped. The time point where chains intersect can be considered as approximate convergence. All the draws prior to this intersection point are treated as a so-called burn-in period and discarded because they are draws from non-equilibrium distributions.

In addition to a graphical test, it is common to do some sort of formal statistical test of convergence. A popular convergence test is Gelman-Rubin's test (Gelman and Rubin, 1992). The basis for the test is a comparison of between-chain variance to within-chain variance, for at least two chains initiated from overdispersed points. From simulation studies, convergence can be concluded when $1 \leq \hat{R} \leq 1.05$, where \hat{R} is the Gelman-Rubin statistic.

While there are many other graphical and numerical convergence diagnostics, they each have to be conducted for each parameter. In practice,

with complex Bayesian models only a few diagnostics are feasible. I propose using a combination of trace plots and Gelman-Rubin tests to examine proposed Bayesian models, and treat chains as converged when both show signatures of convergence. Both methods have been implemented in the R package CODA (Plummer et al., 2006), and can be easily applied to custom Metropolis-Hastings algorithms.

4.2.5 Characterizing Parameter Dependencies

The overall question motivating this study is the existence and nature of correlations among cognitive parameters important for simple decisions. The overall question can be broken up into three questions amenable to statistical analysis: what is the pattern of correlations, how well do models predict different aspects of behavioral data structure and how do their overall predictive accuracies of each model compare when adjusted for complexity. Given samples from the three posterior densities, each question can be addressed in the Bayesian framework.

One way to characterize the correlation structure is to quantify probabilities of correlations in some direction. Using posterior estimates, I propose to estimate probabilities of positive (or negative) correlation for each parameter using the Monte Carlo estimator, which is still appropriate even though samples are dependent (Robert and Casella, 2004; Givens and Hoeting, 2012). For example, the probability that initial position bias and drift rate are pos-

itively correlated is

$$P(\rho^{\beta,\delta} > 0) \approx \frac{1}{N} \sum_{i=1}^N \mathbb{I}_{\{\rho_i^{\beta,\delta} > 0\}}, \quad (4.24)$$

where N is the number of MCMC draws kept after the burn-in. The probability of correlation is a rough measure of evidence for the existence of a correlation structure.

In addition to quantifying probabilities of correlations, we want to understand plausible correlation values suggested by the data. I propose estimating plausible values with a highest posterior density (HPD) interval (Casella and Berger, 2002; Gelman et al., 2013a). An HPD interval is a set of parameter values $\{\theta : f(\theta | \mathbf{x}) \geq k\}$, where k satisfies

$$\int_{\theta: f(\theta|\mathbf{x}) \geq k} f(\theta | \mathbf{x}) d\theta = 1 - \alpha. \quad (4.25)$$

Estimating this interval from a posterior MCMC sample requires solving an optimization problem

$$p^* = \operatorname{argmin} [F^{-1}(1 - \alpha + p) - F^{-1}(p)], \quad (4.26)$$

where p is the lower tail probability, α controls the credibility level and F^{-1} is a quantile function. I will use the recently developed R package SPIn that uses quadratic programming to find p and bootstrapping to minimize Monte Carlo error (Liu et al., 2013). The estimated interval will have the shortest length for a given probability value, and the length will reflect uncertainty about the true correlation value.

Last, I will use the correlation estimates to calculate Kendall's, and tail dependence coefficient for the copula-based models. Both measures will pro-

vide more insight into the correlation structure. The Kendall's coefficient can replace the linear correlation measure by probing for a monotonic relation between processing parameters, and the tail dependence coefficient can provide evidence for extreme combinations of processing parameters. The formulas for both have been specified in Equation 3.27 and Equation 4.10, respectively. I will summarize both quantities with a HPD.

4.2.6 Posterior Predictive Model Comparison

In addition to estimating the nature of correlations, how well models predict various features of the benchmark dataset provides another perspective on presence of dependencies between parameters. In Bayesian analysis, inference about model adequacy is based on examining a posterior predictive density of a given data model. A posterior predictive density is defined as

$$f(\mathbf{y} \mid \mathbf{x}) = \int_{\Theta} f(\mathbf{y} \mid \boldsymbol{\theta}) \pi(\boldsymbol{\theta} \mid \mathbf{x}) d\boldsymbol{\theta} = E_{\boldsymbol{\theta} \mid \mathbf{x}}[f(\mathbf{y} \mid \boldsymbol{\theta})], \quad (4.27)$$

where $\mathbf{y} \in \mathbb{R}^p$ is out-of-sample data, $\mathbf{x} \in \mathbb{R}^p$ is the sample data and $\boldsymbol{\theta} \in \Theta \subseteq \mathbb{R}^m$ are the model parameters. The posterior predictive density describes future or hypothetical replications of observed data, under the same conditions and of the same sample size, after averaging over remaining uncertainty about parameters. The density $f(\mathbf{y} \mid \mathbf{x})$ depends on the data model and hence provides an opportunity to test the model against actual data. In our case, if the correlation structure is a necessary assumption, then a model with an independence copula should show systematic misfit when its predictions are compared to data.

The nature of the comparison between observations and posterior predictive density is formalized by methods of posterior predictive checks (Gelman et al., 2000, 2013a). Because models are simplifications aimed to capture major elements of a data-generating mechanism, checks are conducted with respect to a limited number of data features. The logic of checking a data model is frequentist. Suppose we simulate many replications of the actual experiment according to the posterior predictive density and calculate some statistic of interest. A model is adequate if the statistic based on observed data is likely under the distribution of a statistic based on predicted data.

More concretely, the adequacy of a data model can be inferred based on a statistical test with Bayesian p-values or informally based on visual examination of plots (Gelman et al., 2013a). I will ignore the former and describe a graphical method under a typical assumption that a data model precludes analytical derivation of a posterior predictive density. A graphical procedure requires selecting a discrepancy statistic $s(\cdot)$ that captures some interesting data feature and contrasting observed with predicted values of that statistic. For example, in the case of behavioral data, we could take pairs of mean response times and response probabilities, separated by correct and incorrect responses, and examine observed versus predicted latency-probability graphs.

More generally, we can perform a posterior predictive check by first calculating a discrepancy statistic $s(\mathbf{x}_{obs})$ for observed data. Predicted replications $s(\mathbf{y}_{rep,i})$ can be drawn by simulating from the data model for a given element randomly drawn from a MCMC sample of parameters. Sufficient

numbers of parameters have to be considered to reflect uncertainty about the parameter, so a sample of, say, 250 may be feasible. Then, we can plot the $s(\mathbf{x}_{obs})$ against $s(\mathbf{y}_{rep,i})$. The location of the observed value relative to the high density region will reveal consistency between the data model and observations.

While there a lot of subjective choices to be made with graphical posterior predictive checks, they may be a powerful way to examine the adequacy of a model (Gelman et al., 2013a). In the context of simple choice behavior, the Ratcliff and Rouder (1998) benchmark dataset shows many data features that have been considered of theoretical importance (Vickers, 1979; Luce, 1986), and can be used for defining discrepancy statistics. For now, I do not specify which statistics to use because it will be more principled to use the conclusions from Study A to guide this choice. In the end, I will use graphical posterior predictive checks to analyze the adequacy of generalized sequential sampling models by simulating 50 replicate data sets to calculate the discrepancy statistics and relating them to data. The comparison of the three models' predictions against diagnostic features of the benchmark dataset should reveal any systematic deficiencies and whether they are alleviated by dependent parameters.

4.2.7 Predictive Accuracy of Models

The final question involves the overall predictive accuracy for each model and the selection of the best account of the benchmark dataset. Model

selection, a kind of statistical inference, arises when several competing models are proposed to describe the same data. Ultimately we would want to select the true model, but realistically we are selecting the most adequate model that can guide further research or analysis.

Suppose we have a model set $\{M_1, M_2, \dots, M_k\}$ and we want to select the “best” model. The best model can be obtained by applying some combination of qualitative and quantitative criteria that allow us to compare models (Myung et al., 2008; Shiffrin et al., 2008; Vandekerckhove et al.). In psychology, important qualitative criteria include interpretability of parameters and the plausibility of the explanation a data model provides. For example, parameters of sequential sampling models can be interpreted as important factors affecting a decision process such as bias, caution and rate of evidence accumulation. The models also provide plausible explanations of many effects through reasonable adjustment of parameters across conditions, such as increase in caution under accuracy instructions (Ratcliff and Smith, 2004). Generalizations of sequential sampling models that are used in this study inherit the desirable qualitative attributes, so the primary means of selecting among them are quantitative.

In contrast to qualitative considerations, a common quantitative approach to model selection is based on model generalizability (Gelman et al., 2013b). Generalizability is defined as model’s ability to predict out-of-sample data collected under the same conditions as the data used to estimate the parameters. Typically, generalizability is expressed as a scalar quantity that is calculated for each model to allow ranking of both nested and non-nested

models. A model with higher generalizability can be considered a better approximation of the actual data-generating mechanism, so inference amounts to selecting the model with the highest generalizability value.

The various proposals for calculating out-of-sample predictive accuracy can be decomposed into two parts: an estimate of predictive accuracy based on observed data and a penalty term. The penalty term is a correction for overfitting from using the same data both for parameter estimation and model evaluation. The overfitting of a model to data is driven by model complexity, so the penalty term can also be interpreted as quantifying model complexity. This quantification that may incorporate the number of model parameters and model's functional form (Myung et al., 2008; Gelman et al., 2013b). A popular and easy-to-calculate generalizability criterion used in Bayesian analysis is the deviance information criterion (DIC) (Spiegelhalter et al., 2002). The DIC equals $\overline{D} + p_D$, where

$$\overline{D} = \mathbb{E}_{\boldsymbol{\theta}|\mathbf{x}}[-2 \log f(\mathbf{x} | \boldsymbol{\theta})] \quad (4.28)$$

is an estimate of predictive accuracy for $\mathbf{x} \in \mathbb{R}^p$ and $\boldsymbol{\theta} \in \Theta \subseteq \mathbb{R}^m$, and

$$p_D = \mathbb{E}_{\boldsymbol{\theta}|\mathbf{x}}[-2 \log f(\mathbf{x} | \boldsymbol{\theta})] + 2 \log f(\mathbf{x} | \mathbb{E}_{\boldsymbol{\theta}|\mathbf{x}}[\boldsymbol{\theta}]) + c \quad (4.29)$$

is a correction for complexity, where $c \in \mathbb{R}$ is a constant that cancels out in comparisons. The DIC was developed from information theory considerations of hierarchical Bayesian models with model complexity quantified with effective number of parameters p_D . It is very easy to estimate from a MCMC sample by replacing expected values with Monte Carlo estimates, but the

DIC tends to pick overfitting models by underestimating model complexity (Gelman et al., 2013b).

As opposed to grounding a generalizability criterion in information theory, it is more consistent with the rest of the Bayesian perspective to start with the posterior predictive density. A reasonable desideratum for a generalizability criterion is that it should support selection of a Bayesian model with the most accurate posterior predictive distribution. The Bayesian predictive information criterion (BPIC) was developed by considering maximization of mean posterior predictive density with respect to a set of data models (Ando, 2007, 2011). The derivation shows that

$$BPIC = \bar{D} + 2p_D, \quad (4.30)$$

which is equivalent to doubling the complexity correction in DIC. For a posterior sample obtained by an MCMC sampler, BPIC can be easily estimated with

$$\widehat{BPIC} = -\frac{6}{N} \sum_{i=1}^N \log f(\mathbf{x} \mid \boldsymbol{\theta}_i) + 4 \log f(\mathbf{x} \mid \bar{\boldsymbol{\theta}}), \quad (4.31)$$

where $\bar{\boldsymbol{\theta}}$ is the posterior mean obtained from the MCMC sample. There are two conditions under which the BPIC is valid with hierarchical Bayesian models (Ando, 2011). First, the derivation assumes that the data model is not the true generating distribution, but that it is not far from the true distribution. The second condition is that the marginal posterior distributions are unimodal and not too skewed. The empirical success of sequential sampling models (Usher and McClelland, 2001; Ratcliff and McKoon, 2008; Brown and Heathcote, 2008) and prior hierarchical fits of the which model suggest that

both of these conditions hold (Vandekerckhove et al., 2011; Wiecki et al., 2013). Therefore, I will use BPIC for model selection.

4.2.8 Results

The three questions motivating this study should produce a rich pattern of results. To characterize the nature of parameter dependencies, probabilities and HDIs of correlation parameters will be calculated. The results are quantitative, but can be presented in a graphical form by juxtaposing values for the diffusion models with normal and t copulas. A single plot, with one subplot for each quantity, will succinctly show results.

Results for the comparison of models' predictions against diagnostic features of the data will be more complex. The results are essentially graphical. I will create separate plots for each data feature and condition by overlaying the data discrepancy statistic over the posterior predictive densities. The exact number and nature of the plots is not clear because the discrepancy statistics have not yet been specified.

Lastly, model comparison results are straightforward. I will use a table presenting 3 BPIC values. The preferred model will be the one with lowest BPIC.

4.2.9 Discussion

A priori, it is interesting, but maybe futile to guess what the analyses may result in and what they mean. Instead of speculating, I will state the lim-

itations of study B. One limitation has to do with not checking sensitivity of conclusions to data and parameter models. Similar to Study A, the data model comes from the Ratcliff diffusion model of a simple decision process. Other successful sequential sampling models exist that make qualitatively different assumptions (Ratcliff and Smith, 2004), so the conclusions of Study B would at best be suggestive about other models.

Another part of data models are assumed copulas. Specifically for the elliptical copulas, it is not clear without checking if the estimated correlation structure and the general conclusion about its existence would change with a non-elliptical dependence among parameters. Expanding the set of compared models could address both issues, and may be a direction for further research, but for practical reasons the proposed thesis only considers three models.

While the data model limitation is shared by both Studies A and B, not doing sensitivity analysis of parameter models is unique to Study A. The distributions I proposed to assign to all the unknown parameters is only one possible set of such distributions and may differ considerably across applications and researchers. The apparent subjectivity of choices about prior densities is a common challenge to the Bayesian framework, but is a general feature of statistical analysis. The limitation of Study B is that it does not include comparison of inferences made with alternative prior densities, which would increase the credibility of the conclusions.

Another limitation comes from analyzing a single dataset generated by a two-choice free response discrimination task. Data from many other tasks, such as auditory signal detection or lexical decision, have been analyzed and

other response time procedure, such as responding upon researcher controlled signal, could be applied. It is possible that the correlation structure changes across sensory modalities or under experimental manipulations other than stimulus quality and instructions. Conclusions drawn from the proposed analyses in Study B will be restricted to the design used by Ratcliff and Rouder (1998).

Last, proposed analyses are limited in various ways. The analysis of models' predictions by means of graphical posterior predictive checks sets limitations on the strength of the conclusions. With a graphical approach only qualitative conclusions about the adequacy of the models can be drawn. The model comparison may be improved by considering formal statistical tests of misfit based on the proposed discrepancy statistics (Gelman et al., 2013a). Quantitative posterior predictive checks would provide a more precise statement about which features, if any, the Ratcliff diffusion model misses due to independent parameters.

The model selection analysis also has limitations. The BPIC quantity is computationally convenient, but does not quantify complexity of models in a completely desirable way. Only the effective number of parameters is counted, but the functional forms of data models are not incorporated. More convincing indices of predictive accuracy would rely on Bayes factors or minimum description length, both of which provide a comprehensive penalty for complexity but are hard to calculate in practice (Myung et al., 2008; Cavagnaro et al., 2013; Gelman et al., 2013b).

Chapter 5

Bibliography

- Anderson, J. R. (2007). *How can the human mind occur in the physical universe?* Oxford University Press.
- Anderson, J. R., Bothell, D., Byrne, M. D., Douglass, S., Lebiere, C., and Qin, Y. (2004). An integrated theory of the mind. *Psychological review*, 111(4):1036.
- Ando, T. (2007). Bayesian predictive information criterion for the evaluation of hierarchical bayesian and empirical bayes models. *Biometrika*, 94(2):443–458.
- Ando, T. (2011). Predictive bayesian model selection. *American Journal of Mathematical and Management Sciences*, 31(1-2):13–38.
- Berger, J. O. (1997). *Statistical decision theory and Bayesian analysis*. Springer, ”second” edition.

- Berkes, P., Wood, F., and Pillow, J. W. (2008). Characterizing neural dependencies with copula models. In *NIPS*, pages 129–136.
- Bernardo, J. (2011). Modern bayesian inference: Foundations and objective methods. *Philosophy of Statistics*.
- Berntsen, J., Espelid, T. O., and Genz, A. (1991). An adaptive algorithm for the approximate calculation of multiple integrals. *ACM Transactions on Mathematical Software (TOMS)*, 17(4):437–451.
- Bogacz, R. (2007). Optimal decision-making theories: linking neurobiology with behaviour. *Trends in cognitive sciences*, 11(3):118–125.
- Bogacz, R., Brown, E., Moehlis, J., Holmes, P., and Cohen, J. D. (2006). The physics of optimal decision making: a formal analysis of models of performance in two-alternative forced-choice tasks. *Psychological review*, 113(4):700.
- Brown, J. W. (2014). The tale of the neuroscientists and the computer: why mechanistic theory matters. *Frontiers in Neuroscience*, 8.
- Brown, S. D. and Heathcote, A. (2008). The simplest complete model of choice response time: Linear ballistic accumulation. *Cognitive psychology*, 57(3):153–178.
- Brown, S. D., Marley, A., Donkin, C., and Heathcote, A. (2008). An integrated model of choices and response times in absolute identification. *Psychological Review*, 115(2):396.

- Bussemeyer, J. R. and Diederich, A. (2010). *Cognitive modeling*. Sage.
- Calabresi, P., Picconi, B., Tozzi, A., Ghiglieri, V., and Di Filippo, M. (2014). Direct and indirect pathways of basal ganglia: a critical reappraisal. *Nature neuroscience*, 17(8):1022–1030.
- Camerer, C. (2003). *Behavioral game theory: Experiments in strategic interaction*. Princeton University Press.
- Casella, G. and Berger, R. L. (2002). *Statistical inference*. Duxbury.
- Cassey, P., Heathcote, A., and Brown, S. D. (2014). Brain and behavior in decision-making. *PLoS Comput Biol*, 10(7).
- Cavagnaro, D. R., I, M. J., and Pitt, M. A. (2013). Mathematical modeling. *The Oxford Handbook of Quantitative Methods*, 1:438–453.
- Cho, R. Y., Nystrom, L. E., Brown, E. T., Jones, A. D., Braver, T. S., Holmes, P. J., and Cohen, J. D. (2002). Mechanisms underlying dependencies of performance on stimulus history in a two-alternative forced-choice task. *Cognitive, Affective, & Behavioral Neuroscience*, 2(4):283–299.
- Cole, M. W., Bassett, D. S., Power, J. D., Braver, T. S., Petersen, S. E., and Cole, M. W. (2014). Intrinsic and task-evoked network architectures of the human brain. *Neuron*.
- Collins, A. G. and Frank, M. J. (2013). Cognitive control over learning: creating, clustering, and generalizing task-set structure. *Psychological review*, 120(1):190.

- Coombs, C. H. (1983). *Psychology and mathematics: An essay on theory*. University of Michigan Press Ann Arbor.
- Craigmile, P. F., Peruggia, M., and Van Zandt, T. (2010). Hierarchical bayes models for response time data. *Psychometrika*, 75(4):613–632.
- Czado, C. (2010). Pair-copula constructions of multivariate copulas. In *Copula theory and its applications*, pages 93–109. Springer.
- Donkin, C., Brown, S., Heathcote, A., and Wagenmakers, E. (2011). Diffusion versus linear ballistic accumulation: Different models for response time, same conclusions about psychological mechanisms. *Psychonomic Bulletin & Review*, 55:140–151.
- Dutilh, G., Forstmann, B. U., Vandekerckhove, J., and Wagenmakers, E.-J. (2013). A diffusion model account of age differences in posterror slowing. *Psychology and aging*, 28(1):64.
- Dutilh, G., Vandekerckhove, J., Tuerlinckx, F., and Wagenmakers, E.-J. (2009). A diffusion model decomposition of the practice effect. *Psychonomic Bulletin & Review*, 16(6):1026–1036.
- Edwards, W. (1965). Optimal strategies for seeking information: Models for statistics, choice reaction times, and human information processing. *Journal of Mathematical Psychology*, 2(2):312–329.
- Edwards, W., Lindman, H., and Savage, L. J. (1963). Bayesian statistical inference for psychological research. *Psychological Review*, 70(3):193.

- Gamerman, D. and Lopes, H. F. (2006). *Markov chain Monte Carlo: stochastic simulation for Bayesian inference*. CRC Press.
- Gao, J., Wong-Lin, K., Holmes, P., Simen, P., and Cohen, J. D. (2009). Sequential effects in two-choice reaction time tasks: decomposition and synthesis of mechanisms. *Neural computation*, 21(9):2407–36.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., and Rubin, D. B. (2013a). *Bayesian data analysis*. CRC press.
- Gelman, A., Goegebeur, Y., Tuerlinckx, F., and Van Mechelen, I. (2000). Diagnostic checks for discrete data regression models using posterior predictive simulations. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 49(2):247–268.
- Gelman, A., Hwang, J., and Vehtari, A. (2013b). Understanding predictive information criteria for bayesian models. *Statistics and Computing*, pages 1–20.
- Gelman, A. and Rubin, D. B. (1992). Inference from iterative simulation using multiple sequences. *Statistical science*, pages 457–472.
- Genz, A. C. and Malik, A. (1980). Remarks on algorithm 006: An adaptive algorithm for numerical integration over an n-dimensional rectangular region. *Journal of Computational and Applied Mathematics*, 6(4):295–302.
- Givens, G. H. and Hoeting, J. A. (2012). *Computational statistics*, volume 708. John Wiley & Sons.

- Gold, J. I. and Shadlen, M. N. (2007). The neural basis of decision making. *Annual review of neuroscience*, 30:535–74.
- Goldfarb, S., Wong-Lin, K., Schwemmer, M., Leonard, N. E., and Holmes, P. (2012). Can post-error dynamics explain sequential reaction time patterns? *Frontiers in psychology*, 3.
- Hare, T. A., Schultz, W., Camerer, C. F., O’Doherty, J. P., and Rangel, A. (2011). Transformation of stimulus value signals into motor commands during simple choice. *Proceedings of the National Academy of Sciences*, 108(44):18120–18125.
- Heathcote, A. and Hayes, B. (2012). Diffusion versus linear ballistic accumulation: different models for response time with different conclusions about psychological mechanisms? *Canadian Journal of Experimental Psychology/Revue canadienne de psychologie expérimentale*, 66(2):125.
- Heitz, R. P. (2014). The speed-accuracy tradeoff: history, physiology, methodology, and behavior. *Frontiers in neuroscience*, 8(6):150.
- Hofert, M., Kojadinovic, I., Maechler, M., and Yan, J. (2013). copula: Multivariate dependence with copulas. *R package version 0.999-7*.
- Joe, H. (1997). *Multivariate models and multivariate dependence concepts*, volume 73. CRC Press.
- Johnson, N. L., Kotz, S., and Balakrishnan, N. (1994). *Continuous univariate distributions*, volume 1. John Wiley and Sons.

- Johnson, N. L., Kotz, S., and Balakrishnan, N. (1995). *Continuous univariate distributions*, volume 2. John Wiley and Sons.
- Johnson, S. G. and Narasimhan, B. (2013). *cubature: Adaptive multivariate integration over hypercubes*. R package version 1.1-2.
- Jones, M., Curran, T., Mozer, M. C., and Wilder, M. H. (2013). Sequential effects in response time reveal learning mechanisms and event representations. *Psychological review*, 120(3):628–66.
- Jones, M. and Dzhafarov, E. N. (2014). Unfalsifiability and mutual translatability of major modeling schemes for choice reaction time. *Psychological review*, 121(1):1.
- Jones, M., Kinoshita, S., and Mozer, M. C. (2009). Optimal response initiation: Why recent experience matters. In Koller, D., Schuurmans, D., Bengio, Y., and Bottou, L., editors, *Advances in Neural Information Processing Systems 21*, pages 785–792. Curran Associates, Inc.
- Karlin, S. and Taylor, H. M. (1975). *A first course in stochastic processes*. Academic Press.
- Karlin, S. and Taylor, H. M. (1981). *A second course in stochastic processes*, volume 2. Gulf Professional Publishing.
- Kotz, S., Balakrishnan, N., and Johnson, N. L. (2004). *Continuous multivariate distributions, models and applications*, volume 1. John Wiley and Sons.

- Kruschke, J. K. (2010). What to believe: Bayesian methods for data analysis. *Trends in cognitive sciences*, 14(7):293–300.
- Laming, D. R. J. (1968). Information theory of choice-reaction times.
- Lee, M. D. and Wagenmakers, E.-J. (2014). *Bayesian cognitive modeling: A practical course*. Cambridge University Press.
- Lewandowsky, S. and Farrell, S. (2010). *Computational modeling in cognition: Principles and practice*. Sage.
- Link, S. and Heath, R. (1975). A sequential theory of psychological discrimination. *Psychometrika*, 40(1):77–105.
- Liu, Y., Gelman, A., and Zheng, T. (2013). Simulation-efficient shortest probability intervals. *arXiv preprint arXiv:1302.2142*.
- Ljung, L. (1999). System identification: theory for the user.
- Ljung, L. (2010). Perspectives on system identification. *Annual Reviews in Control*, 34(1):1–12.
- Luce, R. D. (1986). *Response Times: Their Role in Inferring Elementary Mental Organization*. Oxford University Press.
- Macmillan, N. A. and Creelman, C. D. (2004). *Detection theory: A user's guide*. Psychology press.

- Matzke, D. and Wagenmakers, E.-J. (2009). Psychological interpretation of the ex-gaussian and shifted wald parameters: A diffusion model analysis. *Psychonomic Bulletin and Review*, 16(5):798–817.
- Miller, E. K. and Cohen, J. D. (2001). An integrative theory of prefrontal cortex function. *Annual review of neuroscience*, 24:167–202.
- Mozer, M. C., Kinoshita, S., and Shettel, M. (2007). Sequential dependencies in human behavior offer insights into cognitive control. *Integrated models of cognitive systems*, pages 180–193.
- Mulder, M. J., van Maanen, L., and Forstmann, B. U. (2014). Perceptual decision neurosciences—a model-based review. *Neuroscience*.
- Myung, I. J. and Pitt, M. A. (1997). Applying occams razor in modeling cognition: A bayesian approach. *Psychonomic Bulletin & Review*, 4(1):79–95.
- Myung, J. I., Karabatsos, G., and Iverson, G. J. (2008). A statisticians view on bayesian evaluation of informative hypotheses. In *Bayesian evaluation of informative hypotheses*, pages 309–327. Springer.
- Navarro, D. J. and Fuss, I. G. (2009). Fast and accurate calculations for first-passage times in wiener diffusion models. *Journal of Mathematical Psychology*, 53(4):222–230.
- Nelsen, R. B. (2007). *An introduction to copulas*. Springer, 2 edition.

- Peruggia, M., Van Zandt, T., and Chen, M. (2002). Was it a car or a cat i saw? an analysis of response times for word recognition. *LECTURE NOTES IN STATISTICS-NEW YORK-SPRINGER VERLAG*, pages 319–334.
- Peters, E., Hart, P. S., Tusler, M., and Fraenkel, L. (2014). Numbers matter to informed patient choices a randomized design across age and numeracy levels. *Medical Decision Making*, 34(4):430–442.
- Pike, R. (1973). Response latency models for signal detection. *Psychological Review*, 80(1):53.
- Plummer, M., Best, N., Cowles, K., and Vines, K. (2006). CODA: Convergence diagnosis and output analysis for MCMC. *R News*, 6(1):7–11.
- Polanía, R., Krajbich, I., Grueschow, M., and Ruff, C. C. (2014). Neural oscillations and synchronization differentially support evidence accumulation in perceptual and value-based decision making. *Neuron*, 82(3):709–720.
- Polson, N. G., Scott, J. G., et al. (2012). On the half-cauchy prior for a global scale parameter. *Bayesian Analysis*, 7(4):887–902.
- R Core Team (2014). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.
- Rabbitt, P. (1979). How old and young subjects monitor and control responses for accuracy and speed. *British Journal of Psychology*, 70(2):305–311.

- Rae, B., Heathcote, A., Donkin, C., Averell, L., and Brown, S. (2014). The hare and the tortoise: emphasizing speed can change the evidence used to make decisions. *Journal of Experimental Psychology: Learning, Memory, and Cognition*.
- Ratcliff, R. (1978). A theory of memory retrieval. *Psychological review*, 85(2):59.
- Ratcliff, R. (1993). Methods for dealing with reaction time outliers. *Psychological bulletin*, 114(3):510.
- Ratcliff, R. (2002). A diffusion model account of response time and accuracy in a brightness discrimination task: Fitting real data and failing to fit fake but plausible data. *Psychonomic Bulletin & Review*, 9(2):278–291.
- Ratcliff, R. (2013). Parameter variability and distributional assumptions in the diffusion model. *Psychological review*, 120(1):281.
- Ratcliff, R. (2014). Measuring psychometric functions with the diffusion model.
- Ratcliff, R., Hasegawa, Y. T., Hasegawa, R. P., Smith, P. L., and Segraves, M. A. (2007). Dual diffusion model for single-cell recording data from the superior colliculus in a brightness-discrimination task. *Journal of neurophysiology*, 97(2):1756–1774.
- Ratcliff, R., Love, J., Thompson, C. A., and Opfer, J. E. (2012). Children are

- not like older adults: A diffusion model analysis of developmental changes in speeded responses. *Child development*, 83(1):367–381.
- Ratcliff, R. and McKoon, G. (2008). The diffusion decision model: Theory and data for two-choice decision tasks. *Neural computation*, 20(4):873–922.
- Ratcliff, R. and Rouder, J. N. (1998). Modeling response times for two-choice decisions. *Psychological Science*, 9(5):347–356.
- Ratcliff, R. and Smith, P. L. (2004). A comparison of sequential sampling models for two-choice reaction time. *Psychological review*, 111(2):333.
- Ratcliff, R. and Tuerlinckx, F. (2002). Estimating parameters of the diffusion model: Approaches to dealing with contaminant reaction times and parameter variability. *Psychonomic bulletin & review*, 9(3):438–481.
- Ratcliff, R., Van Zandt, T., and McKoon, G. (1999). Connectionist and diffusion models of reaction time. *Psychological review*, 106(2):261.
- Rinkenauer, G., Osman, A., Ulrich, R., Müller-Gethmann, H., and Mattes, S. (2004). On the locus of speed-accuracy trade-off in reaction time: inferences from the lateralized readiness potential. *Journal of Experimental Psychology: General*, 133(2):261.
- Robert, C. P. and Casella, G. (2004). *Monte Carlo statistical methods*, volume 319. Citeseer.
- Ross, S. M. (2014). *Introduction to probability models*, volume 12. Academic press.

- Rouder, J. N. and Lu, J. (2005). An introduction to bayesian hierarchical models with an application in the theory of signal detection. *Psychonomic Bulletin and Review*, 12(4):573–604.
- Rouder, J. N., Lu, J., Speckman, P., Sun, D., and Jiang, Y. (2005). A hierarchical model for estimating response time distributions. *Psychonomic Bulletin & Review*, 12(2):195–223.
- Rouder, J. N., Morey, R. D., and S, P. M. (2014). *Bayesian hierarchical models*, volume 1.
- Shadlen, M. N. and Kiani, R. (2013). Decision making as a window on cognition. *Neuron*, 80(3):791–806.
- Shadlen, M. N. and Newsome, W. T. (1996). Motion perception: seeing and deciding. *Proceedings of the National Academy of Sciences*, 93(2):628–633.
- Shiffrin, R. M., Lee, M. D., Kim, W., and Wagenmakers, E.-J. (2008). A survey of model evaluation approaches with a tutorial on hierarchical bayesian methods. *Cognitive Science*, 32(8):1248–1284.
- Simon, D. (2013). *Evolutionary optimization algorithms*. John Wiley & Sons.
- Sklar, M. (1959). *Fonctions de répartition à n dimensions et leurs marges*. Université Paris 8.
- Smith, P. L. (1995). Psychophysically principled models of visual simple reaction time. *Psychological review*, 102(3):567.

- Smith, P. L. (2000). Stochastic dynamic models of response time and accuracy: A foundational primer. *Journal of Mathematical Psychology*, 44(3):408–463.
- Smith, P. L. and Ratcliff, R. (2004). Psychology and neurobiology of simple decisions. *Trends in neurosciences*, 27(3):161–168.
- Smith, P. L. and Ratcliff, R. (2009). An integrated theory of attention and decision making in visual signal detection. *Psychological review*, 116(2):283.
- Spiegelhalter, D. J., Best, N. G., Carlin, B. P., and Van Der Linde, A. (2002). Bayesian measures of model complexity and fit. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 64(4):583–639.
- Stone, M. (1960). Models for choice-reaction time. *Psychometrika*, 25(3):251–260.
- Storn, R. and Price, K. (1997). Differential evolution—a simple and efficient heuristic for global optimization over continuous spaces. *Journal of global optimization*, 11(4):341–359.
- Swensson, R. G. (1972). The elusive tradeoff: Speed vs accuracy in visual discrimination tasks. *Perception & Psychophysics*, 12(1):16–32.
- Teodorescu, A. R. and Usher, M. (2013). Disentangling decision models: From independence to competition. *Psychological review*, 120(1):1.

- Ter Braak, C. J. (2006). A markov chain monte carlo version of the genetic algorithm differential evolution: easy bayesian computing for real parameter spaces. *Statistics and Computing*, 16(3):239–249.
- Townsend, J. T. and Ashby, F. G. (1983). *The stochastic modeling of elementary psychological processes*. CUP Archive.
- Tuerlinckx, F. (2004). The efficient computation of the cumulative distribution and probability density functions in the diffusion model. *Behavior Research Methods, Instruments, & Computers*, 36(4):702–716.
- Tuerlinckx, F., Maris, E., Ratcliff, R., and De Boeck, P. (2001). A comparison of four methods for simulating the diffusion process. *Behavior Research Methods, Instruments, & Computers*, 33(4):443–456.
- Turner, B. M., Forstmann, B. U., Wagenmakers, E.-J., Brown, S. D., Sederberg, P. B., and Steyvers, M. (2013a). A bayesian framework for simultaneously modeling neural and behavioral data. *NeuroImage*, 72:193–206.
- Turner, B. M., Sederberg, P. B., Brown, S. D., and Steyvers, M. (2013b). A method for efficiently sampling from distributions with correlated dimensions. *Psychological methods*, 18(3):368.
- Tversky, A. and Kahneman, D. (1981). The framing of decisions and the psychology of choice. *Science*, 211(4481):453–458.
- Usher, M. and McClelland, J. L. (2001). The time course of perceptual choice: the leaky, competing accumulator model. *Psychological review*, 108(3):550.

- van Schouwenburg, M., Aarts, E., and Cools, R. (2010). Dopaminergic modulation of cognitive control: distinct roles for the prefrontal cortex and the basal ganglia. *Current pharmaceutical design*, 16(18):2026–2032.
- Van Zandt, T., Colonius, H., and Proctor, R. W. (2000). A comparison of two response time models applied to perceptual matching. *Psychonomic Bulletin & Review*, 7(2):208–256.
- Van Zandt, T. and Townsend, J. T. (2013). Designs for and analyses of response time experiments. *The Oxford Handbook of Quantitative Methods, Volume 1*, 1:260.
- Vandekerckhove, J., Matzke, D., and Wagenmakers, E.-J. Model comparison and the principle of parsimony. *Oxford Handbook of Computational and Mathematical Psychology*. Oxford University Press, Oxford (to appear).
- Vandekerckhove, J. and Tuerlinckx, F. (2007). Fitting the ratcliff diffusion model to experimental data. *Psychonomic Bulletin & Review*, 14(6):1011–1026.
- Vandekerckhove, J. and Tuerlinckx, F. (2008). Diffusion model analysis with matlab: A dmat primer. *Behavior Research Methods*, 40(1):61–72.
- Vandekerckhove, J., Tuerlinckx, F., and Lee, M. D. (2011). Hierarchical diffusion models for two-choice response times. *Psychological Methods*, 16(1):44.

- Verdonck, S. and Tuerlinckx, F. (2014). The ising decision maker: a binary stochastic network for choice response time. *Psychological review*, 121(3):422.
- Vickers, D. (1979). *Decision processes in visual perception*. Academic Press New York.
- Voss, A., Rothermund, K., and Voss, J. (2004). Interpreting the parameters of the diffusion model: An empirical validation. *Memory & Cognition*, 32(7):1206–1220.
- Voss, A. and Voss, J. (2007). Fast-dm: A free program for efficient diffusion model analysis. *Behavior Research Methods*, 39(4):767–775.
- Wagenmakers, E.-J. (2009). Methodological and empirical developments for the ratcliff diffusion model of response times and accuracy. *European Journal of Cognitive Psychology*, 21(5):641–671.
- Wagenmakers, E.-J., Farrell, S., and Ratcliff, R. (2004). Estimation and interpretation of $1/\alpha$ noise in human cognition. *Psychonomic bulletin & review*, 11(4):579–615.
- Wald, A. and Wolfowitz, J. (1948). Optimum character of the sequential probability ratio test. *The Annals of Mathematical Statistics*, pages 326–339.
- Wang, X.-J. (2008). Decision making in recurrent neuronal circuits. *Neuron*, 60(2):215–234.

Wiecki, T. V., Sofer, I., and Frank, M. J. (2013). Hddm: hierarchical bayesian estimation of the drift-diffusion model in python. *Frontiers in neuroinformatics*, 7.

Zhang, J. and Rowe, J. B. (2014). Dissociable mechanisms of speed-accuracy tradeoff during visual perceptual learning are revealed by a hierarchical drift-diffusion model. *Frontiers in neuroscience*, 8(April):69.